System Busses / Networks-on-Chip

EECE 579 - Advanced Topics in VLSI Design

Spring 2009

Brad Quinton

Outline

1. **Simple systems busses**
   - Overview
   - AMBA APB
   - Advantages/Limitations

2. **Complex systems busses**
   - Overview
   - AMBA AHB
   - Advantages/Limitations

3. **Networks-on-Chip (NoC)**
   - Overview
   - AMBA AXI
   - Research Topics: Topology, Protocol, VLSI Implementation...
   - Review: "A Generic Architecture for On-Chip Packet-Switched Interconnections"

Bluetooth "Platform" SoC

Processor

Application Specific Logic

ARBITER

DECODER

ARM7TDMI

TIC

SMC

Memory Controller

RADIO I/F

ADC

AHB

BRIDGE

APB

SHARED MEMORY CONTROLLER

SPEECH I/F

DMA

POWER & CLOCK CONTROL

PLL CLOCKS

System Bus / Hardware I/F

SHARED MEMORY

LMC

DAP I/F

WATCH DOG

GPIO

PIC

TIMERS

UART

UART

ACI USB

Low-speed I/O and Support Logic

3

Simple System Busses

Simple System Busses

- The primary goal of a simple system bus is to allow **software** (running on a processor) to **communicate** with other **hardware** in the SoC

- There are many different implementation ... but they are all **very similar**

Embedded Processor I/O

- RISC-based embedded processors communicate with external hardware using two simple instructions:
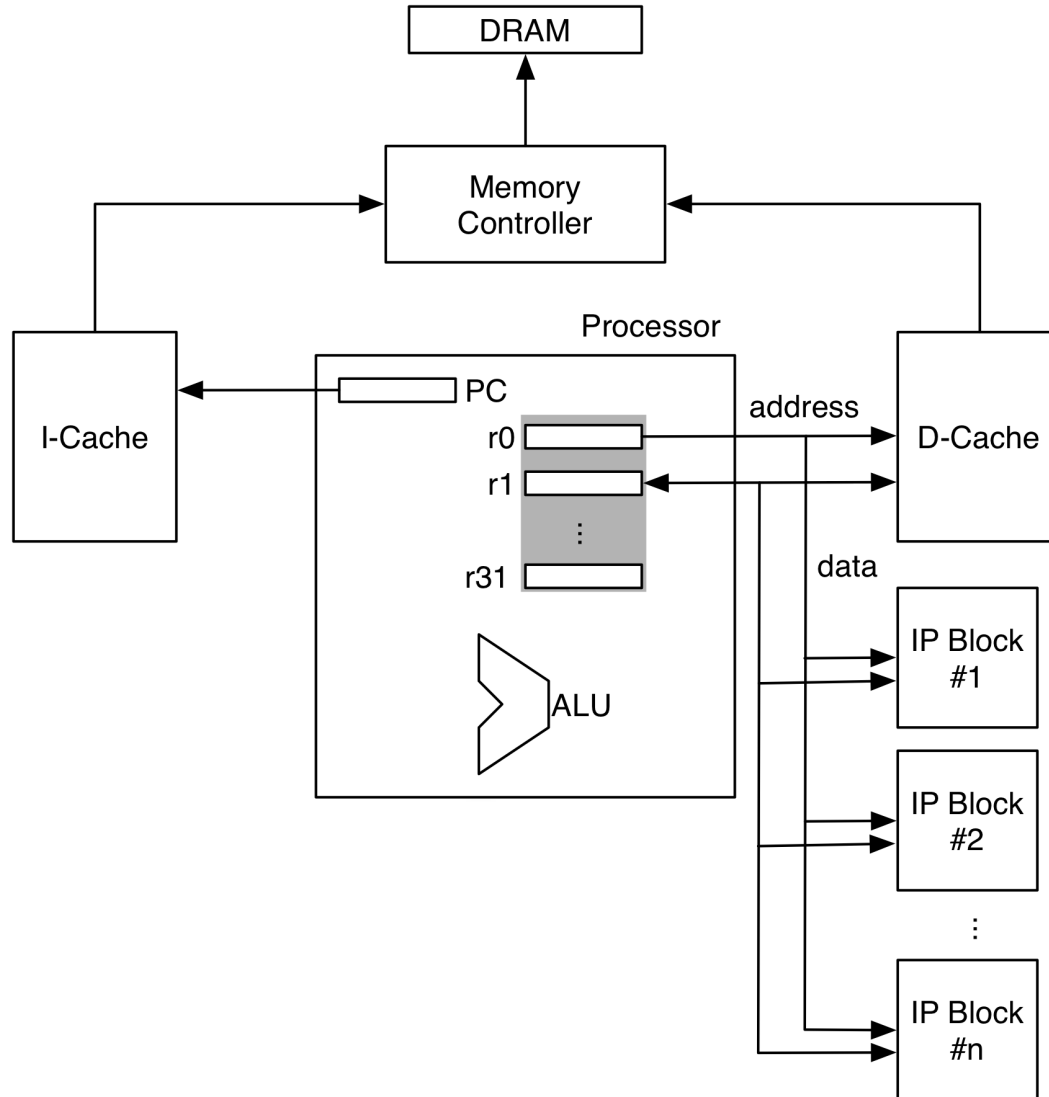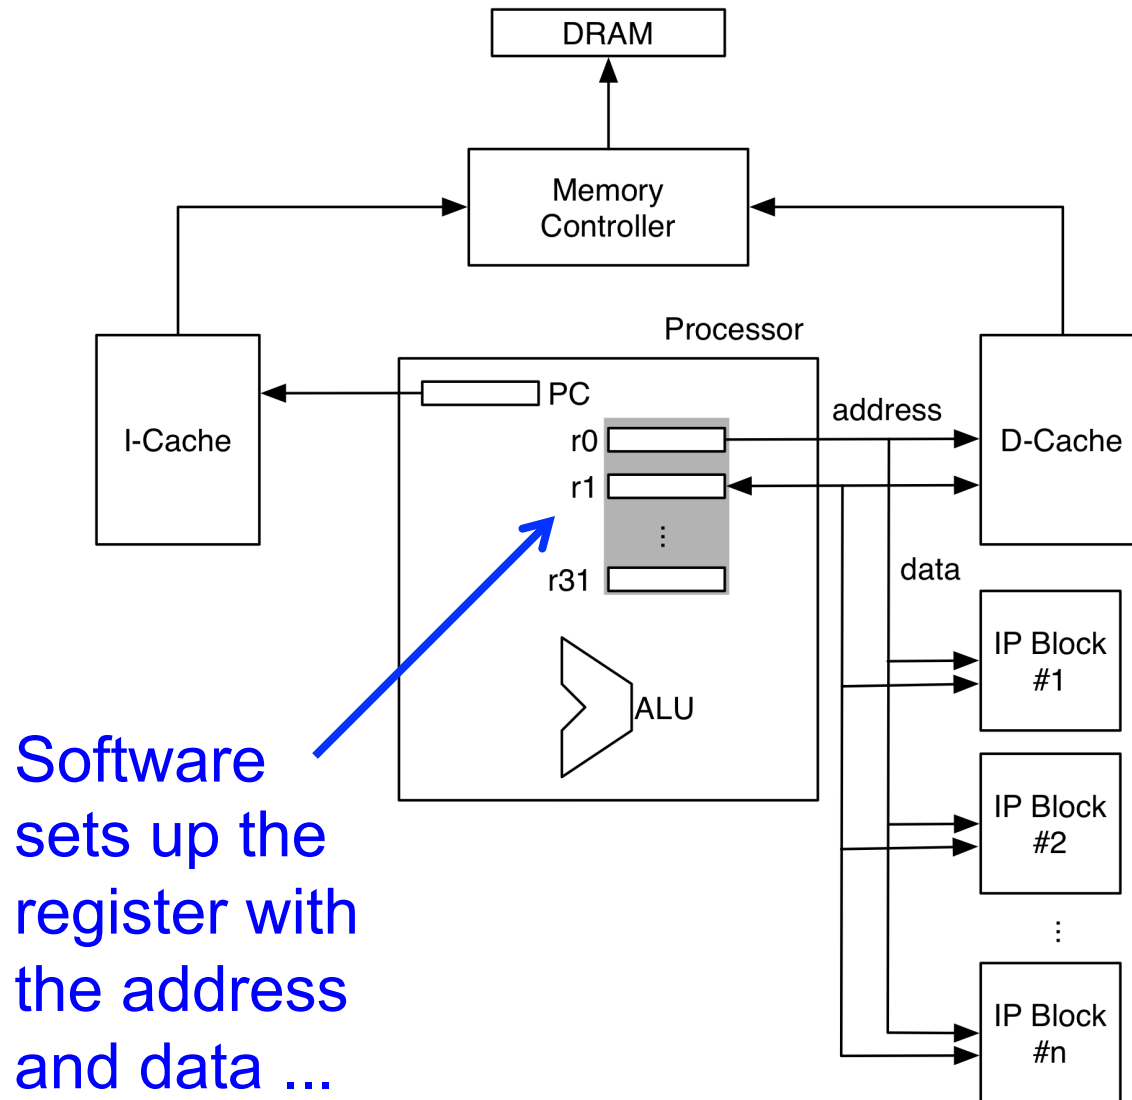
Embedded Processor I/O

- RISC-based embedded processors communicate with external hardware using two simple instructions:

    - **Load Operation:** Copies a word of data from a specific address to a *local register*

    - **Store Operation:** Copies a word of data from a *local register* to a specific address

Embedded Processor I/O

- RISC-based embedded processors communicate with external hardware using two simple instructions:

  – **Load Operation:** Copies a word of data from a specific address to a *local register*

  – **Store Operation:** Copies a word of data from a *local register* to a specific address

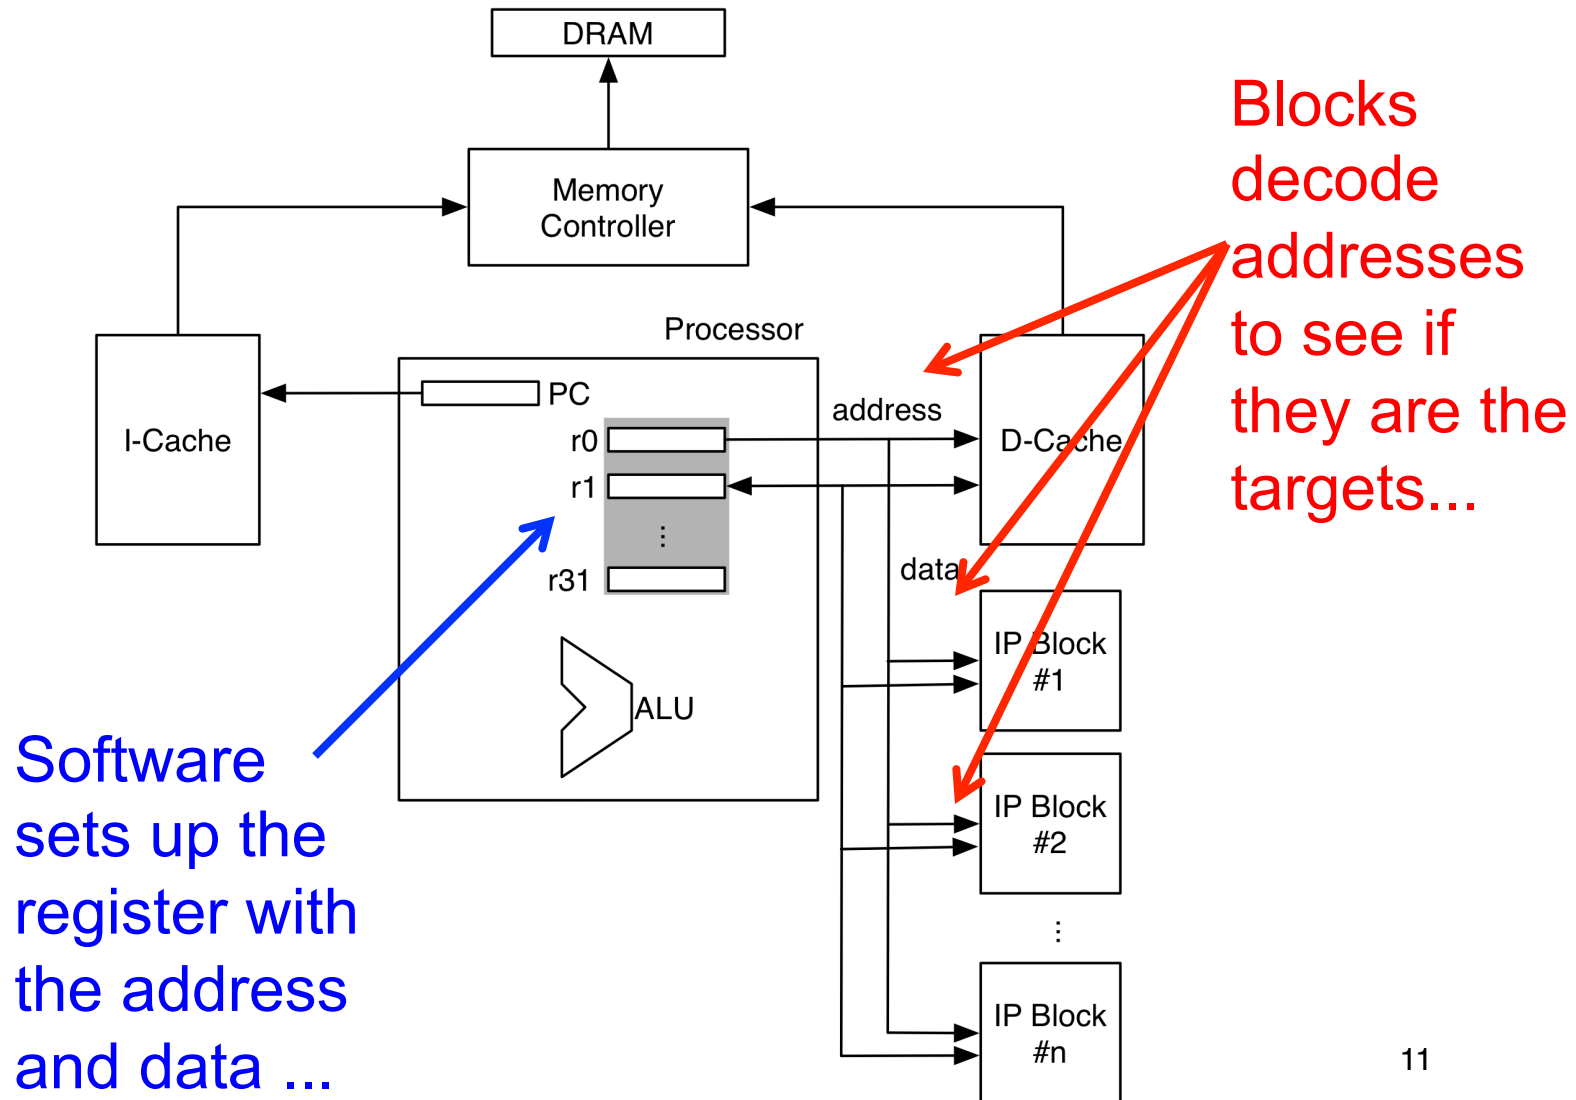- The simple system bus is just a direct extension of this model
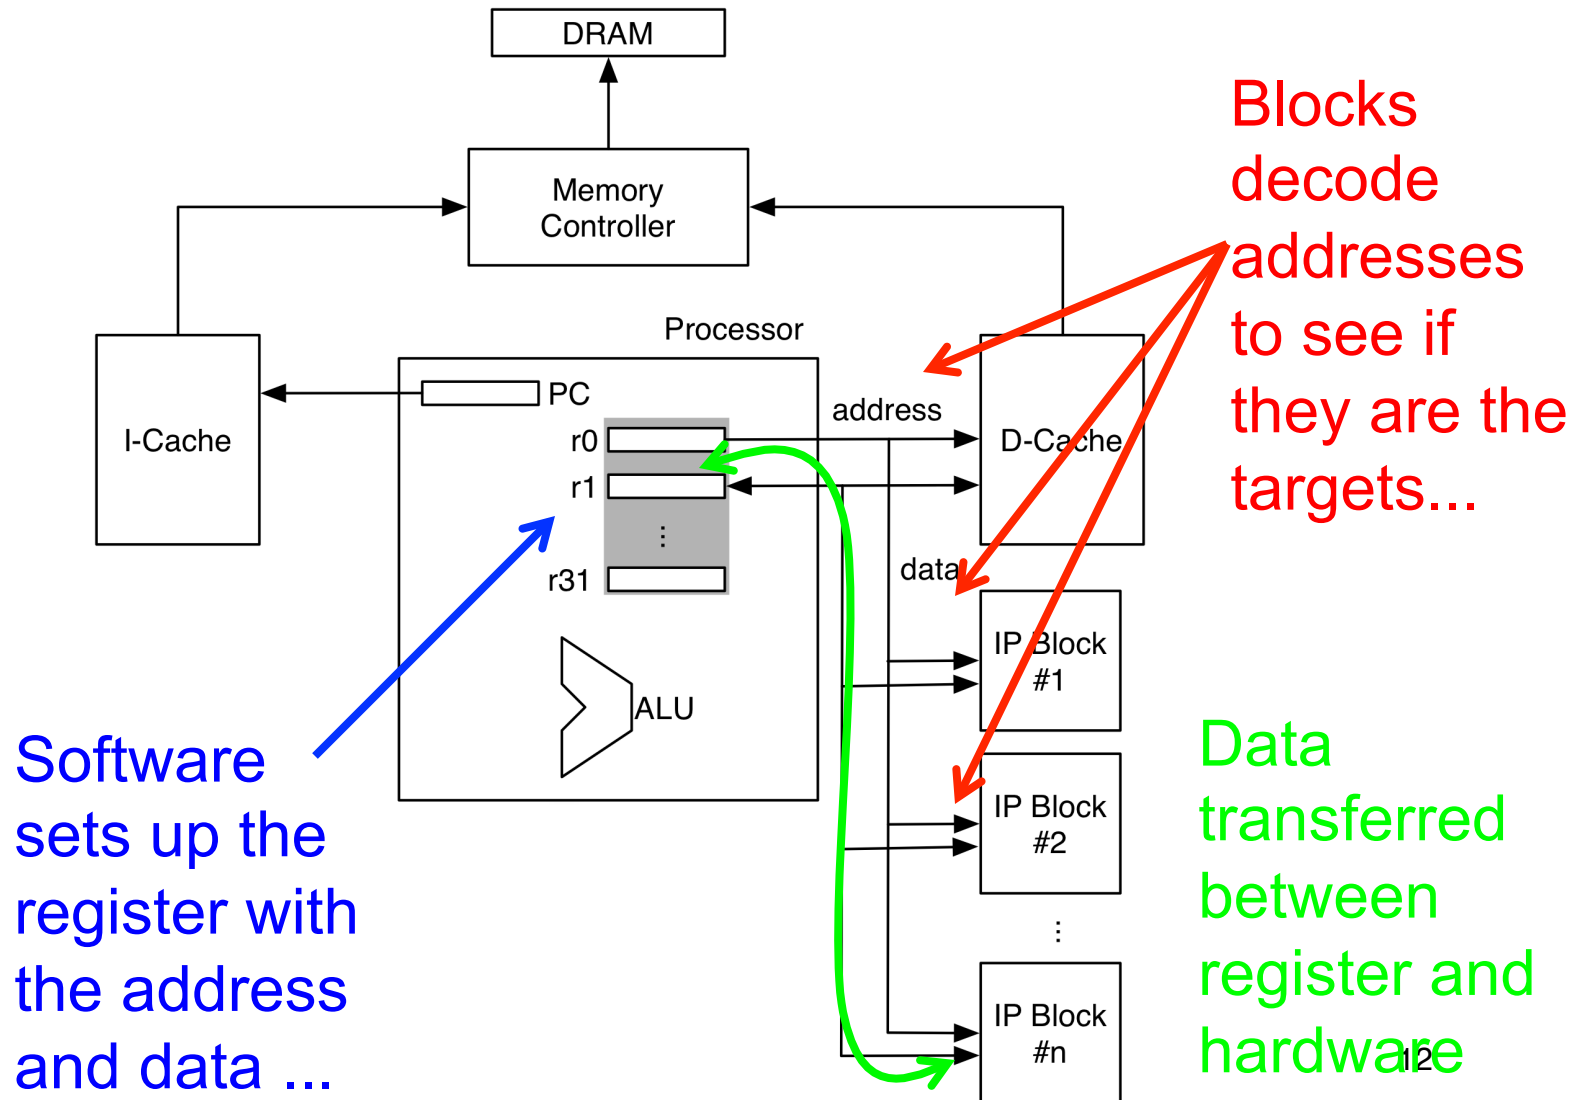
Embedded Processor I/O

Embedded Processor I/O

Software sets up the register with the address and data ...

Embedded Processor I/O

DRAM

Memory Controller

Processor

I-Cache

PC

r0
r1
⋮
r31

ALU

address

data

D-Cache

IP Block #1

IP Block #2

⋮

IP Block #n

Blocks decode addresses to see if they are the targets...

Software sets up the register with the address and data ...

11

Embedded Processor I/O

Blocks decode addresses to see if they are the targets...

Software sets up the register with the address and data ...

Data transferred between register and hardware

DRAM

Memory Controller

Processor

I-Cache

PC

r0
r1
⋮
r31

ALU

address

D-Cache

data

IP Block #1

IP Block #2

⋮

IP Block #n

12

AMBA Specification

- **AMBA**: Advanced Microcontroller Bus Architecture

- Created by ARM to enable standardized interfaces to their embedded processors

- Actually three standards: APB, AHB, and AXI

- Very commonly used for commercial IP cores

AMBA Specification

- **AMBA**: Advanced Microcontroller Bus Architecture

- Created by ARM to enable standardized interfaces to their embedded processors

  Simple Bus

- Actually three standards APB, AHB, and AXI

- Very commonly used for commercial IP cores

AMBA Specification

- **AMBA**: Advanced Microcontroller Bus Architecture

- Created by ARM to enable standardized interfaces to their embedded processors

  Simple Bus   Complex Bus

- Actually three standards APB, AHB, and AXI

- Very commonly used for commercial IP cores

AMBA Specification

- **AMBA**: Advanced Microcontroller Bus Architecture

- Created by ARM to enable standardized interfaces to their embedded processors

Simple Bus   Complex Bus   NoC

- Actually three standards APB, AHB, and AXI

- Very commonly used for commercial IP cores

AMBA APB: Read Operation
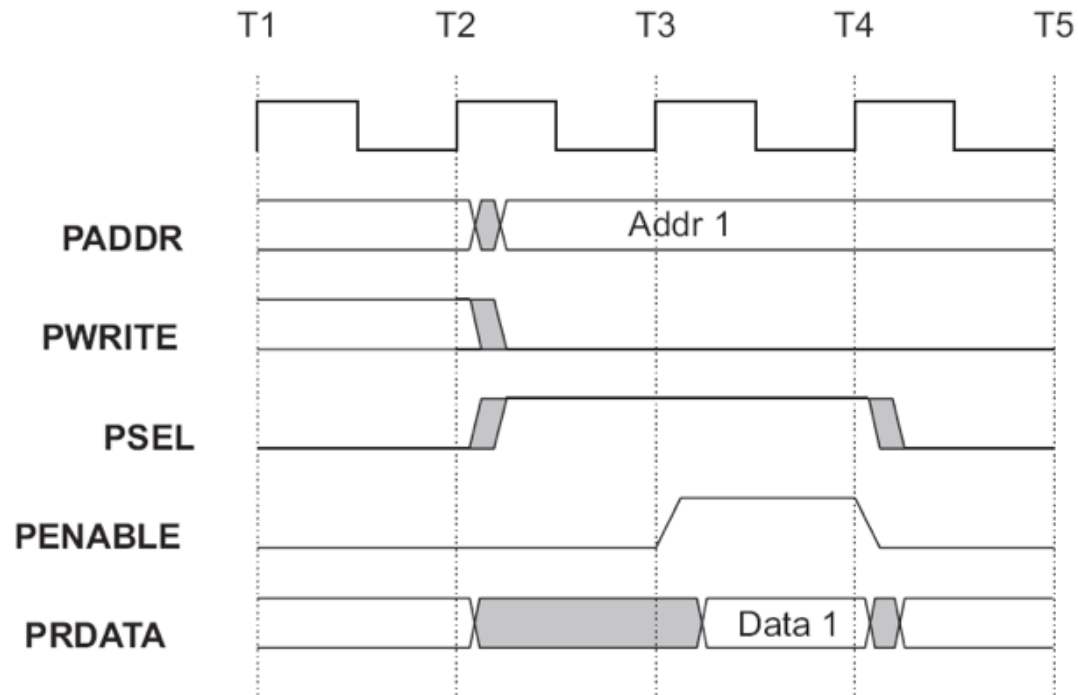
AMBA APB: Read Operation
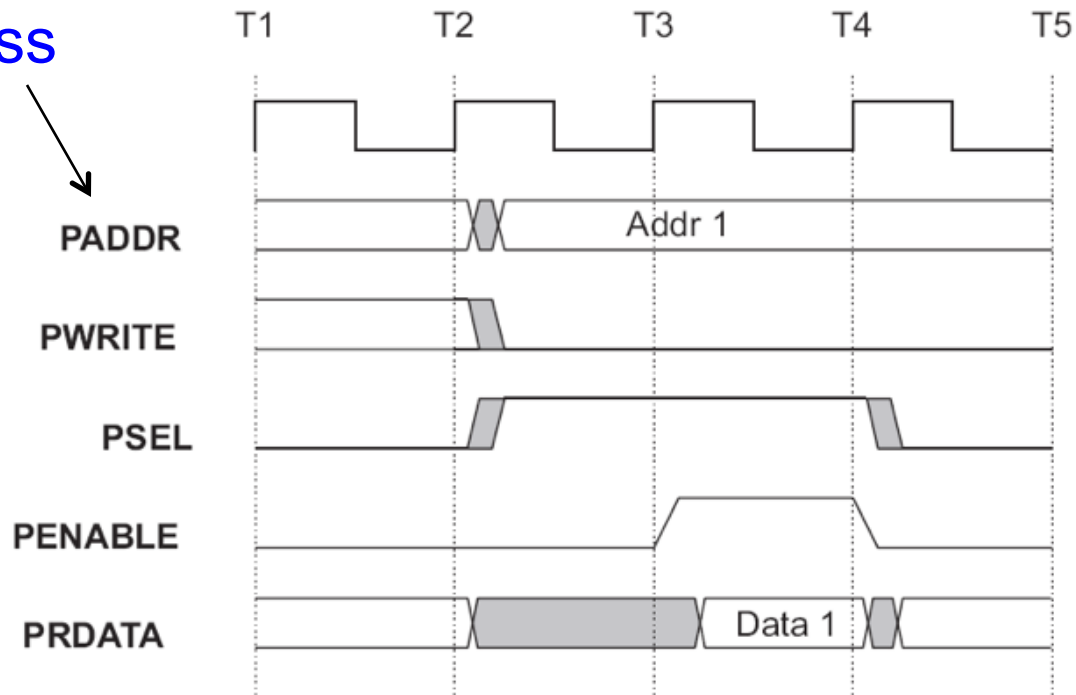
Target Address

AMBA APB: Read Operation

Target Address
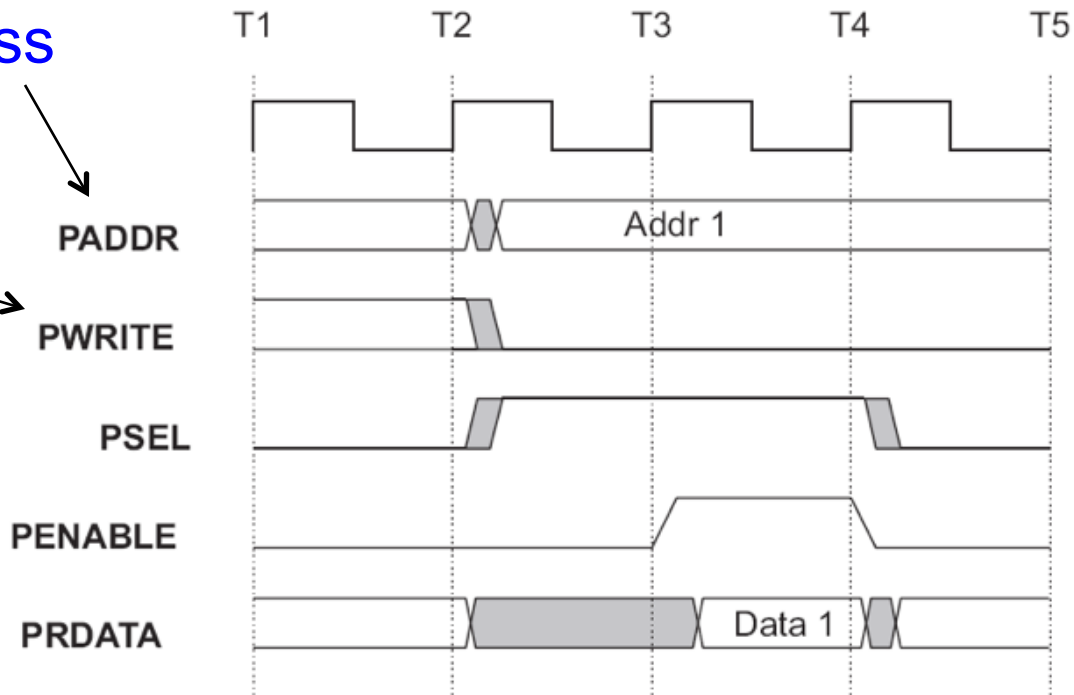
Transaction Type

AMBA APB: Read Operation
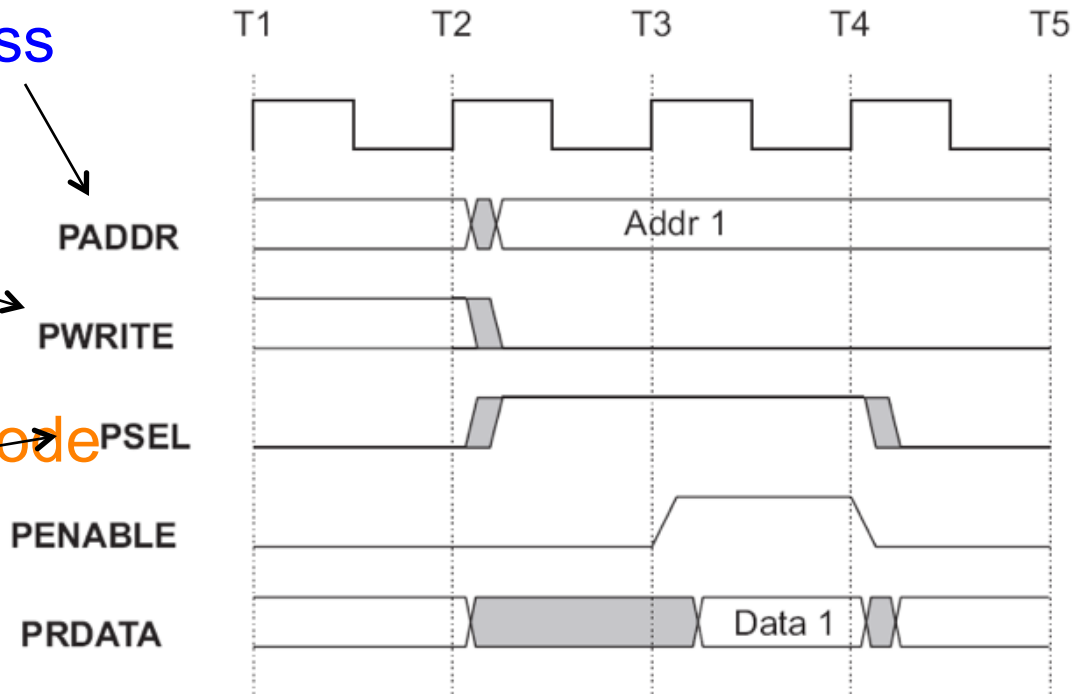
Target Address

Transaction Type

Address Decode

AMBA APB: Read Operation

Target Address

Transaction Type

Address Decode

Optional (for asynchronous implementations ...)

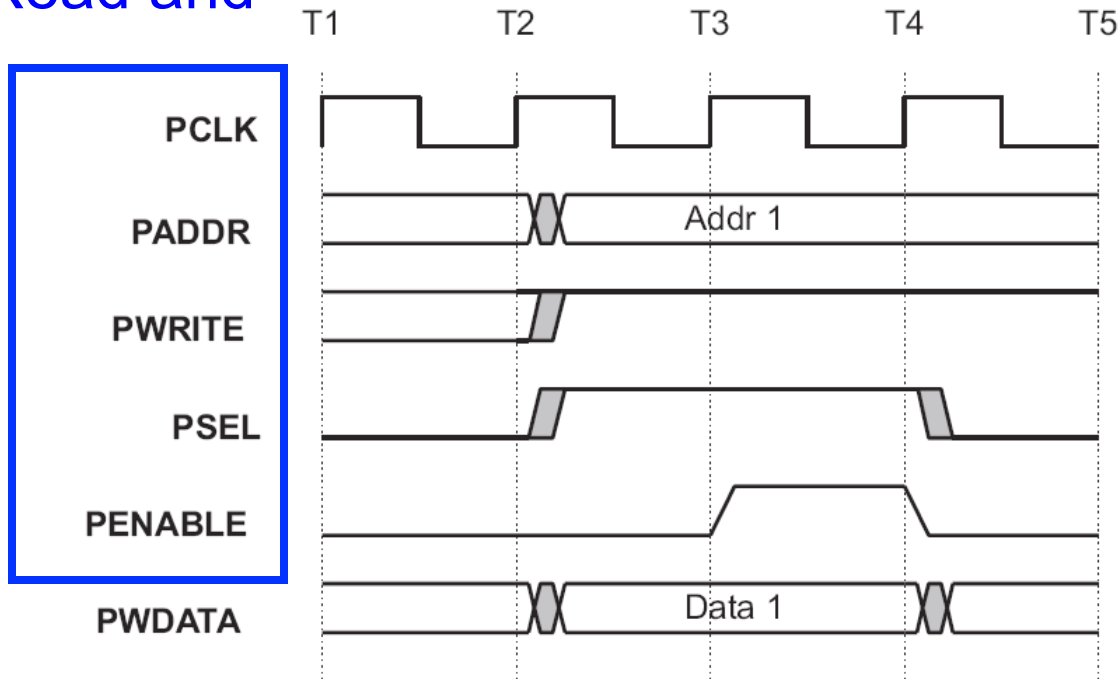# AMBA APB: Read Operation

Target Address

Transaction
Type

Address Decode

Optional (for
asynchronous
implementations
...)

Read Data

T1 T2 T3 T4 T5

PADDR — Addr 1

PWRITE

PSEL

PENABLE

PRDATA — Data 1
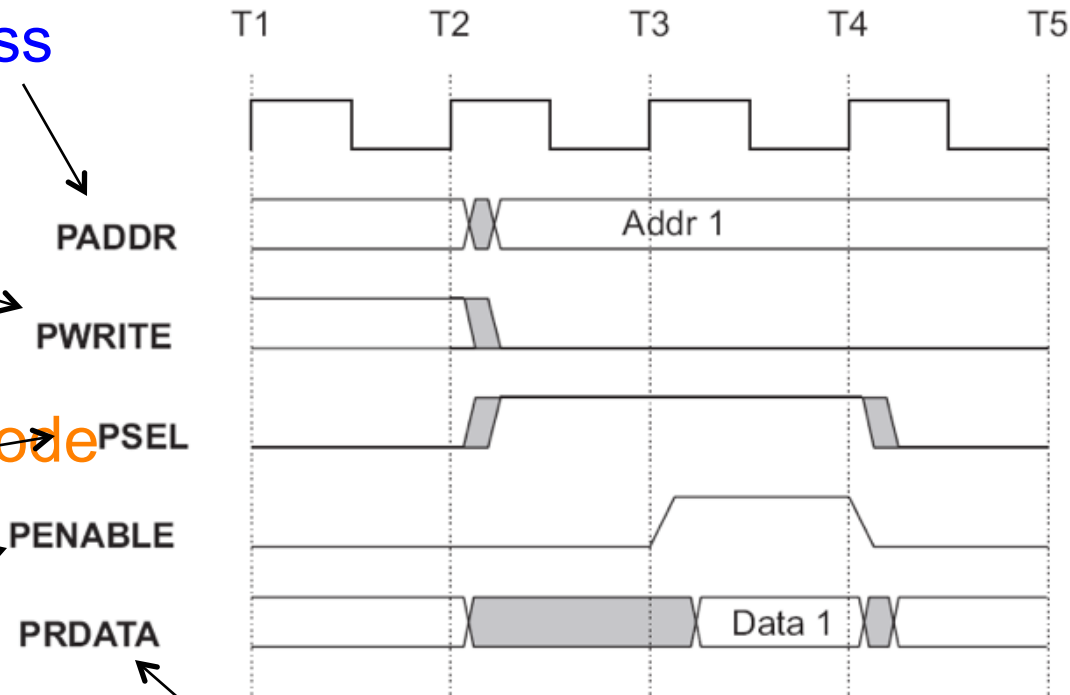
AMBA APB: Write Operation

AMBA APB: Write Operation

Common Signals Between Read and Write

AMBA APB: Write Operation

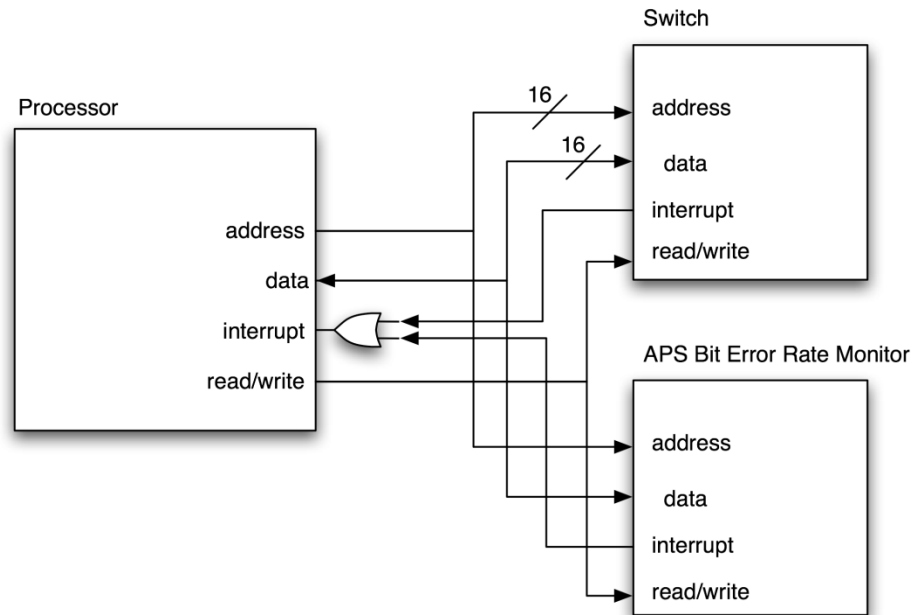Common Signals Between Read and Write

Write Data

Remember Our Case Study

Simple generic processor interface:

- data width:16 bits

- address width: 16 bits

- read cycle time: 50 ns

- write cycle time: 50 ns

Remember Our Case Study

Simple generic processor interface:

- data width:16 bits

- address width: 16 bits

- read cycle time: 50 ns

- write cycle time: 50 ns

Simple Bus Advantages

- Simple to implement
- Easy to understand
- Simple programming model
- Easy to add new hardware blocks
- Minimal hardware requirements (most of the signals are shared)

Simple Bus Limitations

- Single Master - limits parallelism
- Scalability - performance suffers as bus is loaded...
- Single outstanding request - poor throughput and multi-threading performance bottleneck

Case Study: Single Master
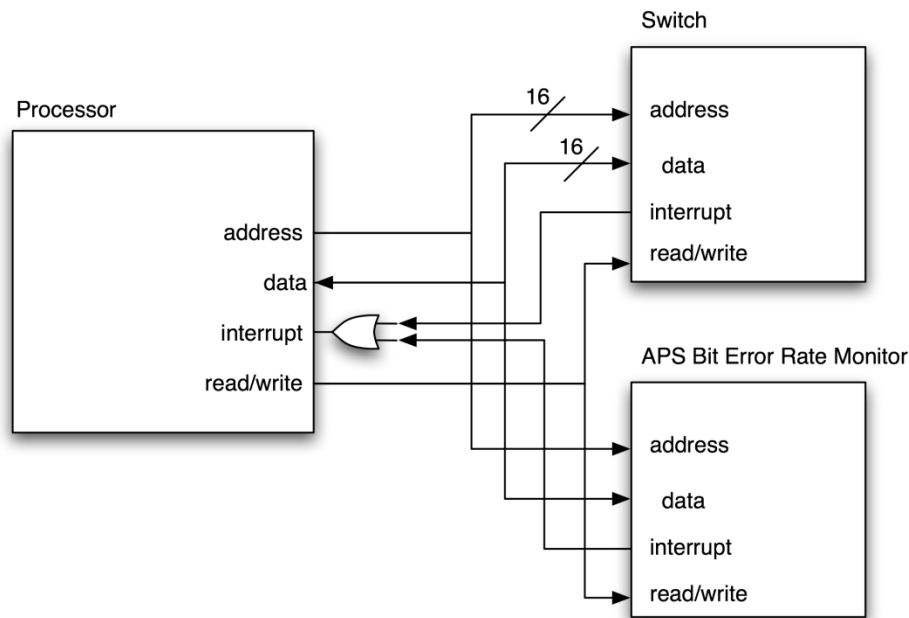
- Imagine a new partition:
  - *APS Bit Error Monitor* communicates directly with *Switch*

- Simple bus doesn't work...

Case Study: Single Master

- Imagine a new partition:
  - *APS Bit Error Monitor* communicates directly with *Switch*

- Simple bus doesn't work...

Case Study: Single Master

No Path

Processor

Switch

16 — address
16 — data
interrupt
read/write

address
data
interrupt
read/write

APS Bit Error Rate Monitor

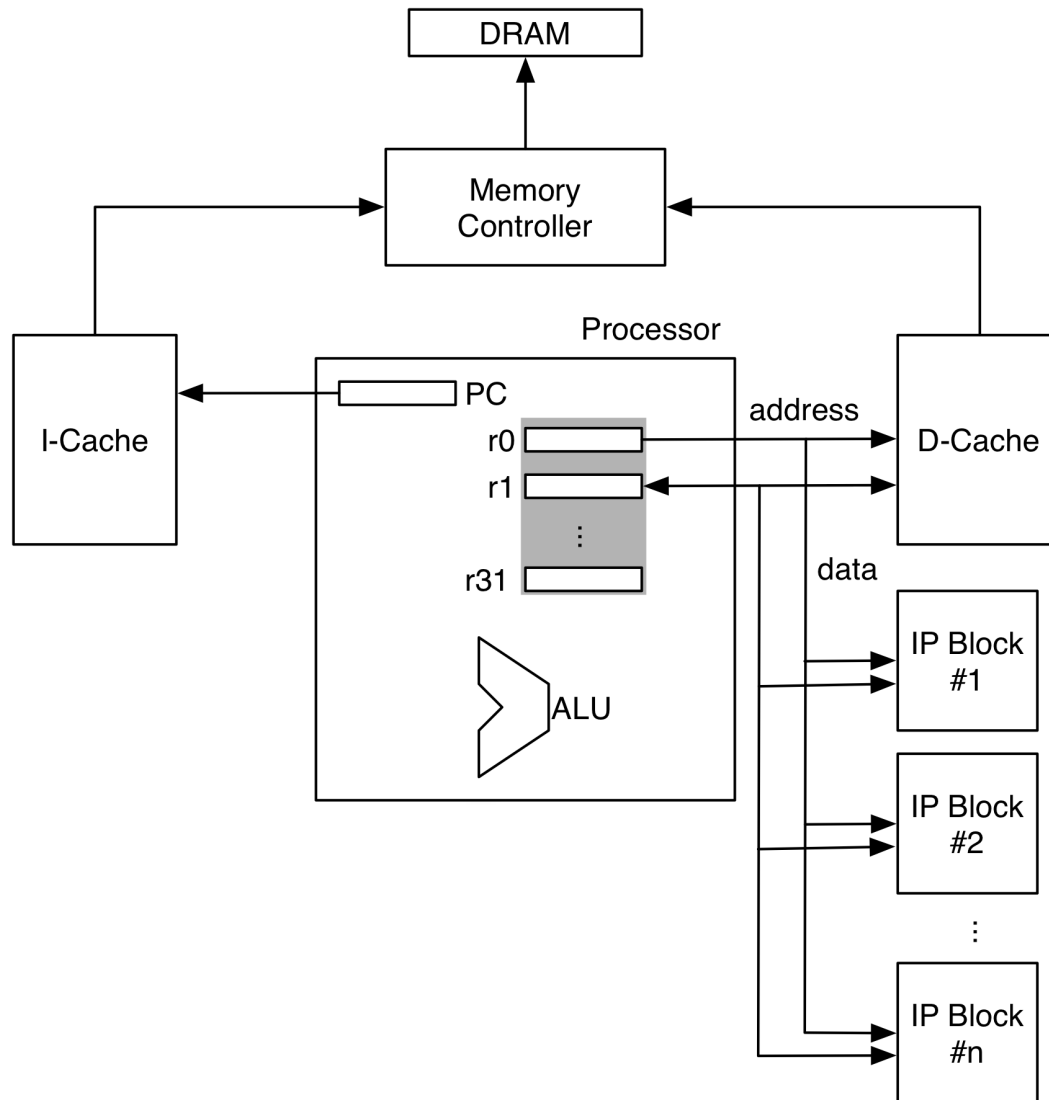address
data
interrupt
read/write

- Imagine a new partition:
  - *APS Bit Error Monitor* communicates directly with *Switch*

- Simple bus doesn't work...

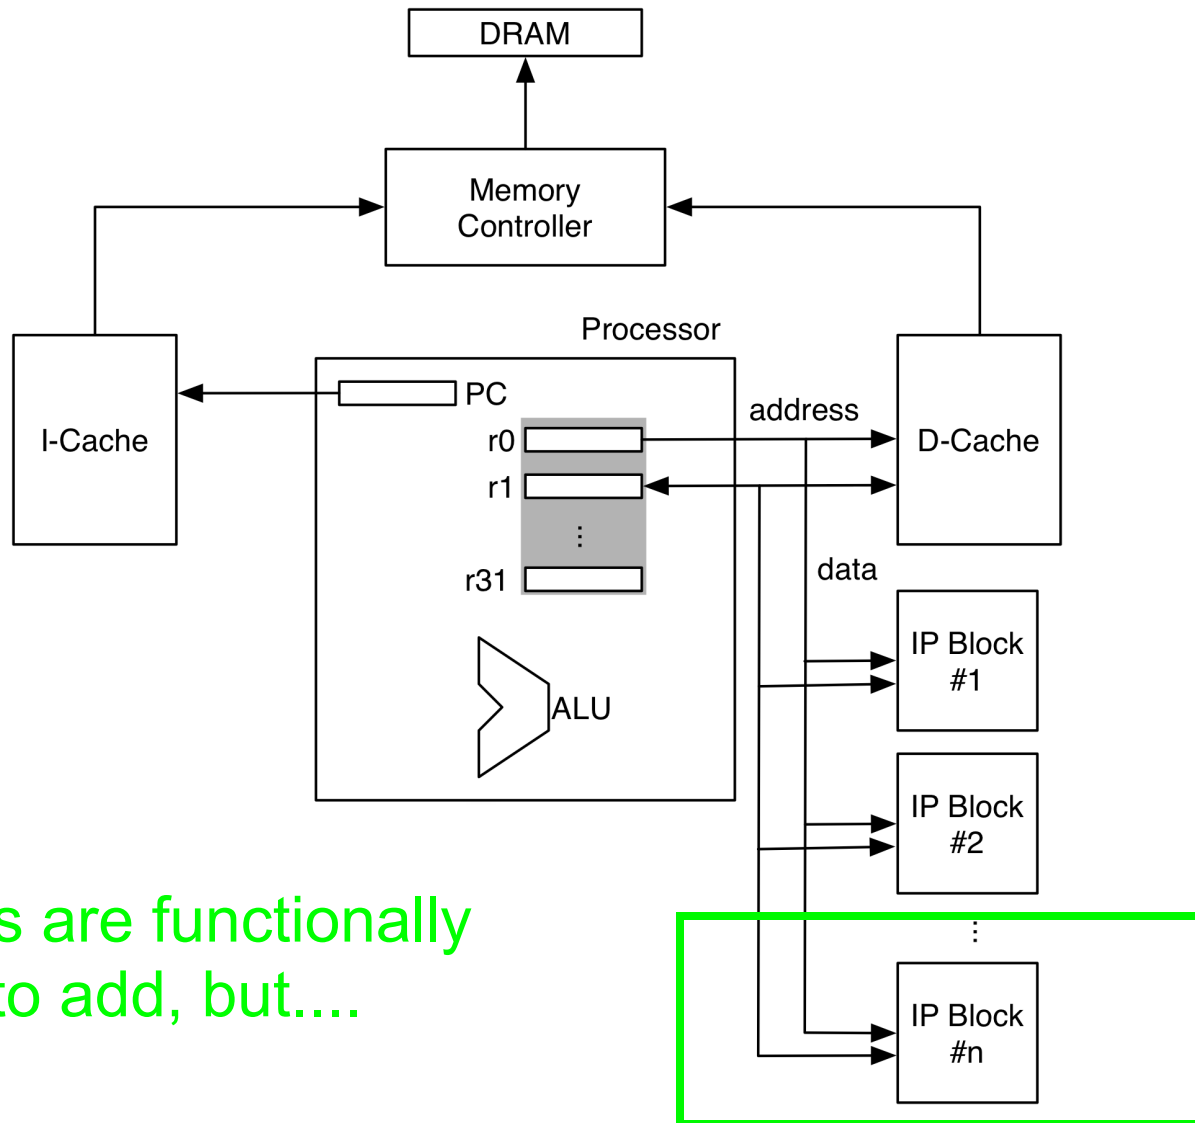- This can make software the bottleneck in the system....

Single Master Summary

- A bus that is limited to a single master:

  - Makes inter-block communication inefficient
  - Limits parallelism between hardware and software
  - Increases reliance on interrupts
  - Creates software performance bottlenecks
  - Is not compatible with multiple processors
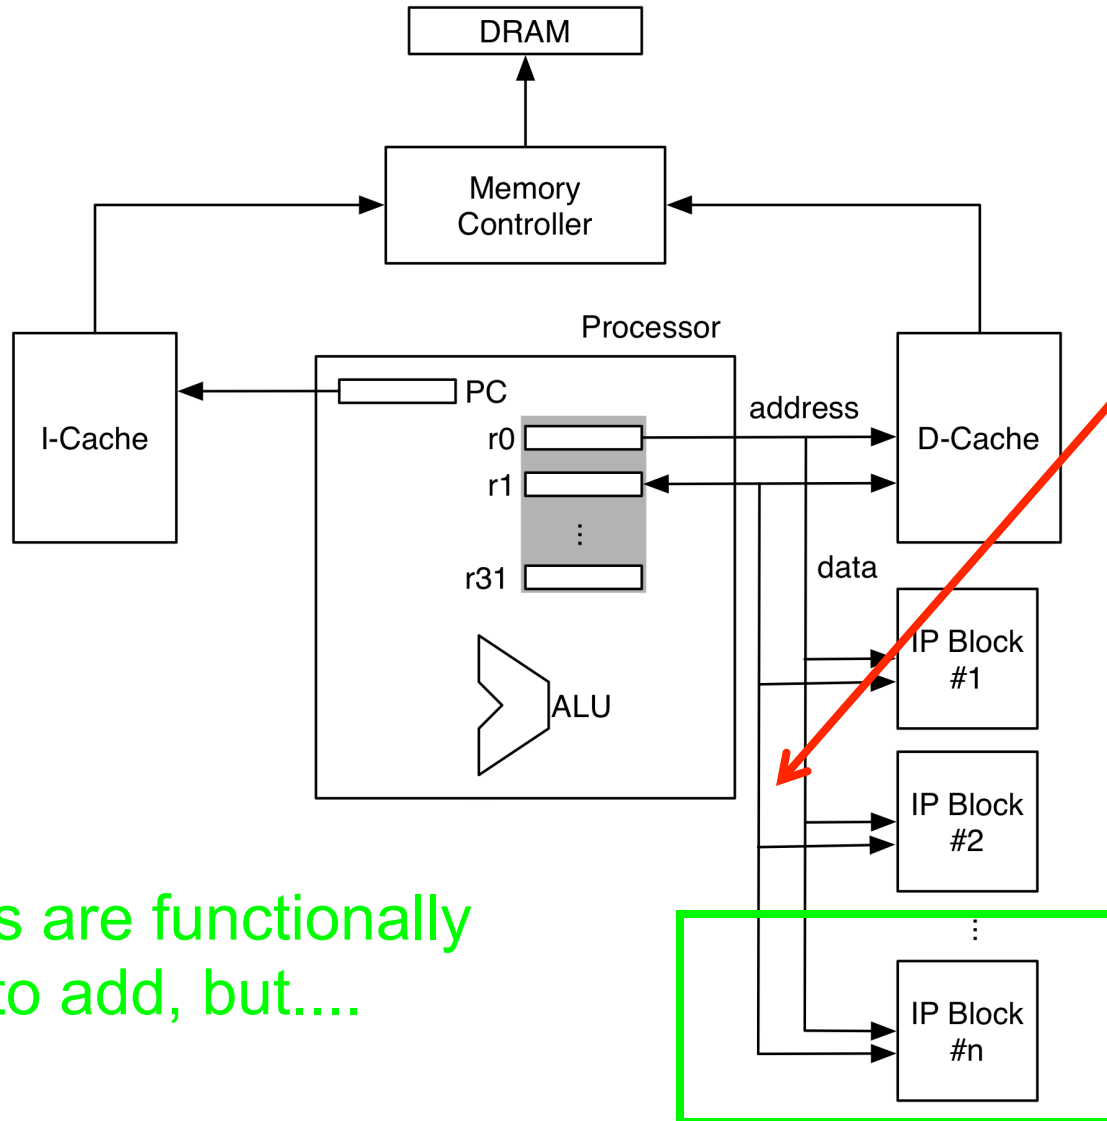
Scalability

Scalability

DRAM

Memory Controller

Processor

I-Cache

PC

r0
r1
⋮
r31

address

ALU

D-Cache

data

IP Block #1

IP Block #2

⋮

IP Block #n

Blocks are functionally easy to add, but....

Scalability

Each new block increases the delay on the address and data

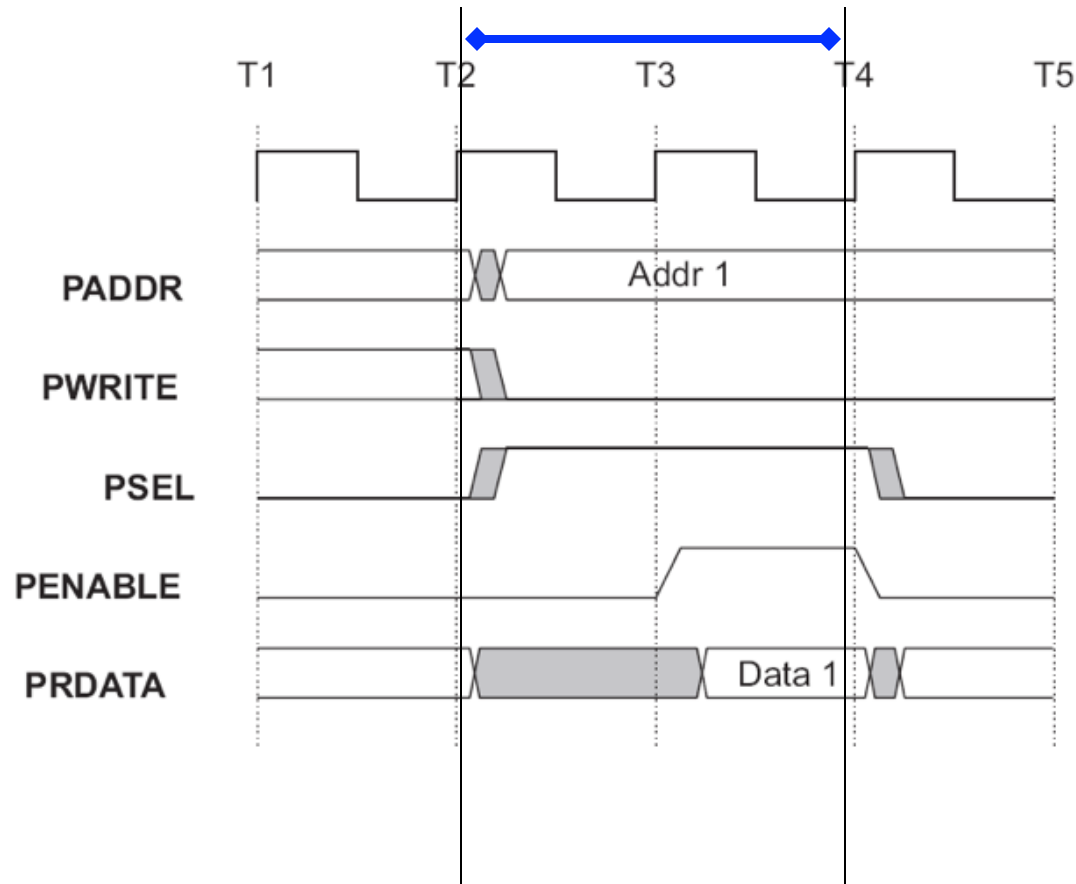Blocks are functionally easy to add, but....

Scalability Summary

- Simple busses are not scaleable because:

    – The address and data "fan-out" to each target
    – Adding a new block increases the load on the bus
    – Increased fanout + greater load = reduce performance

Single Outstanding Request

Single Outstanding Request

Processor is stalled waiting for response...

Single Outstanding Request

Processor is stalled waiting for response...

best-case <= 50% efficiency

Single Outstanding Request Summary

- Busses limited to a single outstanding request:

  – Reduce software performance since the software must "stall" on the first transaction

  – Are not able to achieve full bus throughput since the *data bus* is idle during the *address phase*

Complex System Busses

Complex Systems Busses

- The complex system bus is attempts to address some of the issues with the simple bus:

    - Multi-master
    - Pipelined transactions

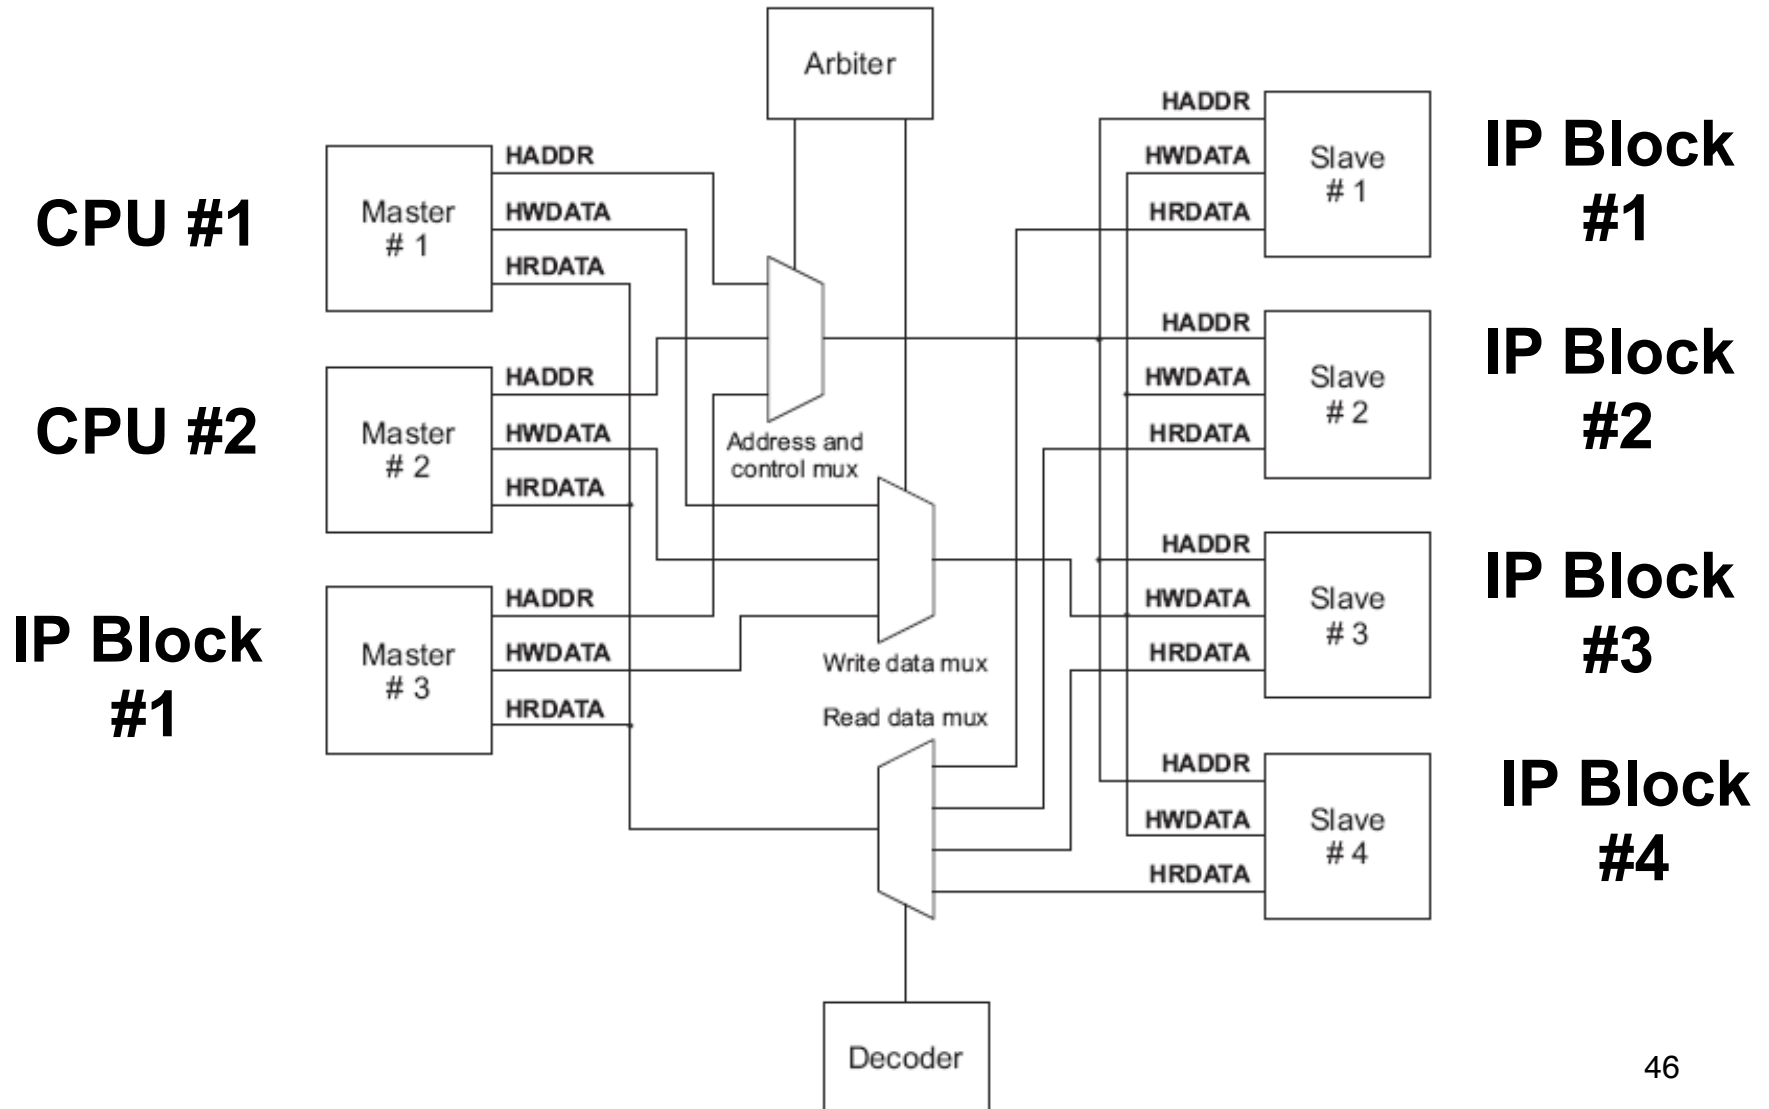- There are many different ways to go about this...

AMBA AHB

- AHB addresses many of the limitations of APB:

  - multi-master
  - multiple outstanding transactions (sort of...)
  - back-to-back transactions

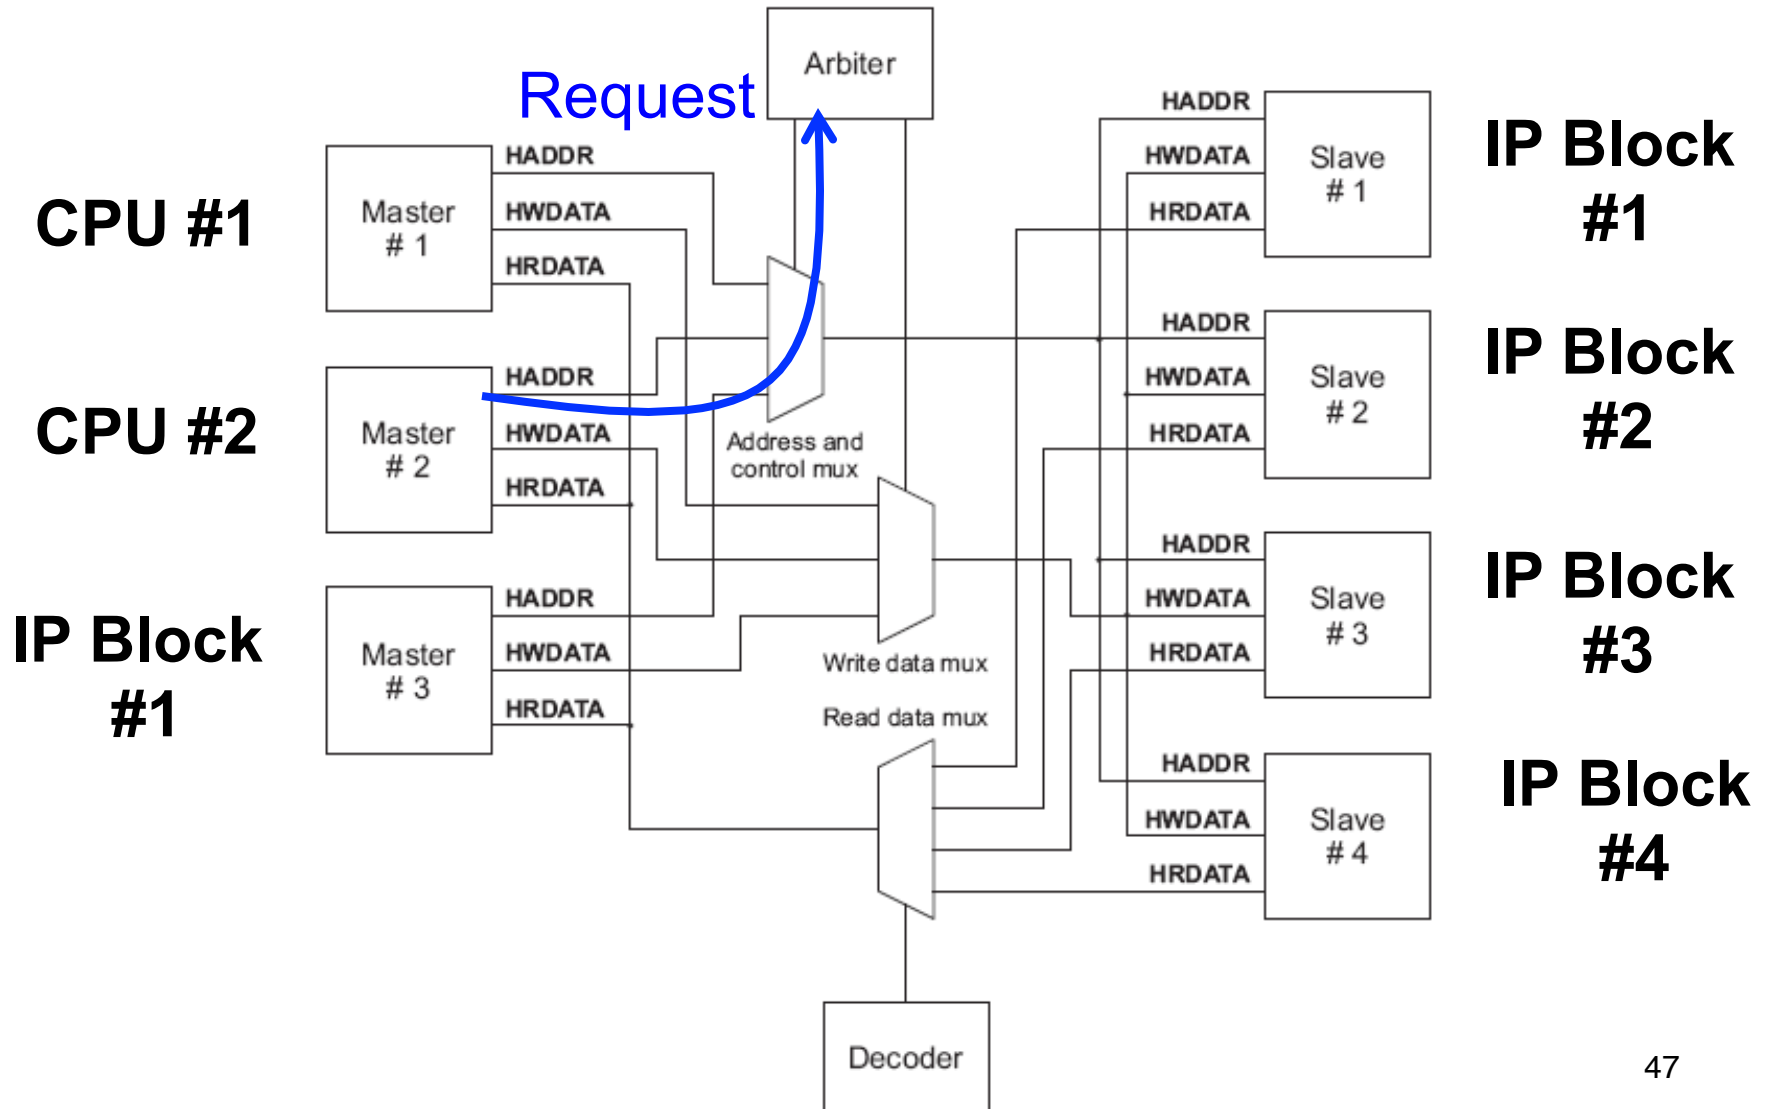- Unfortunately, this adds **significant complexity**
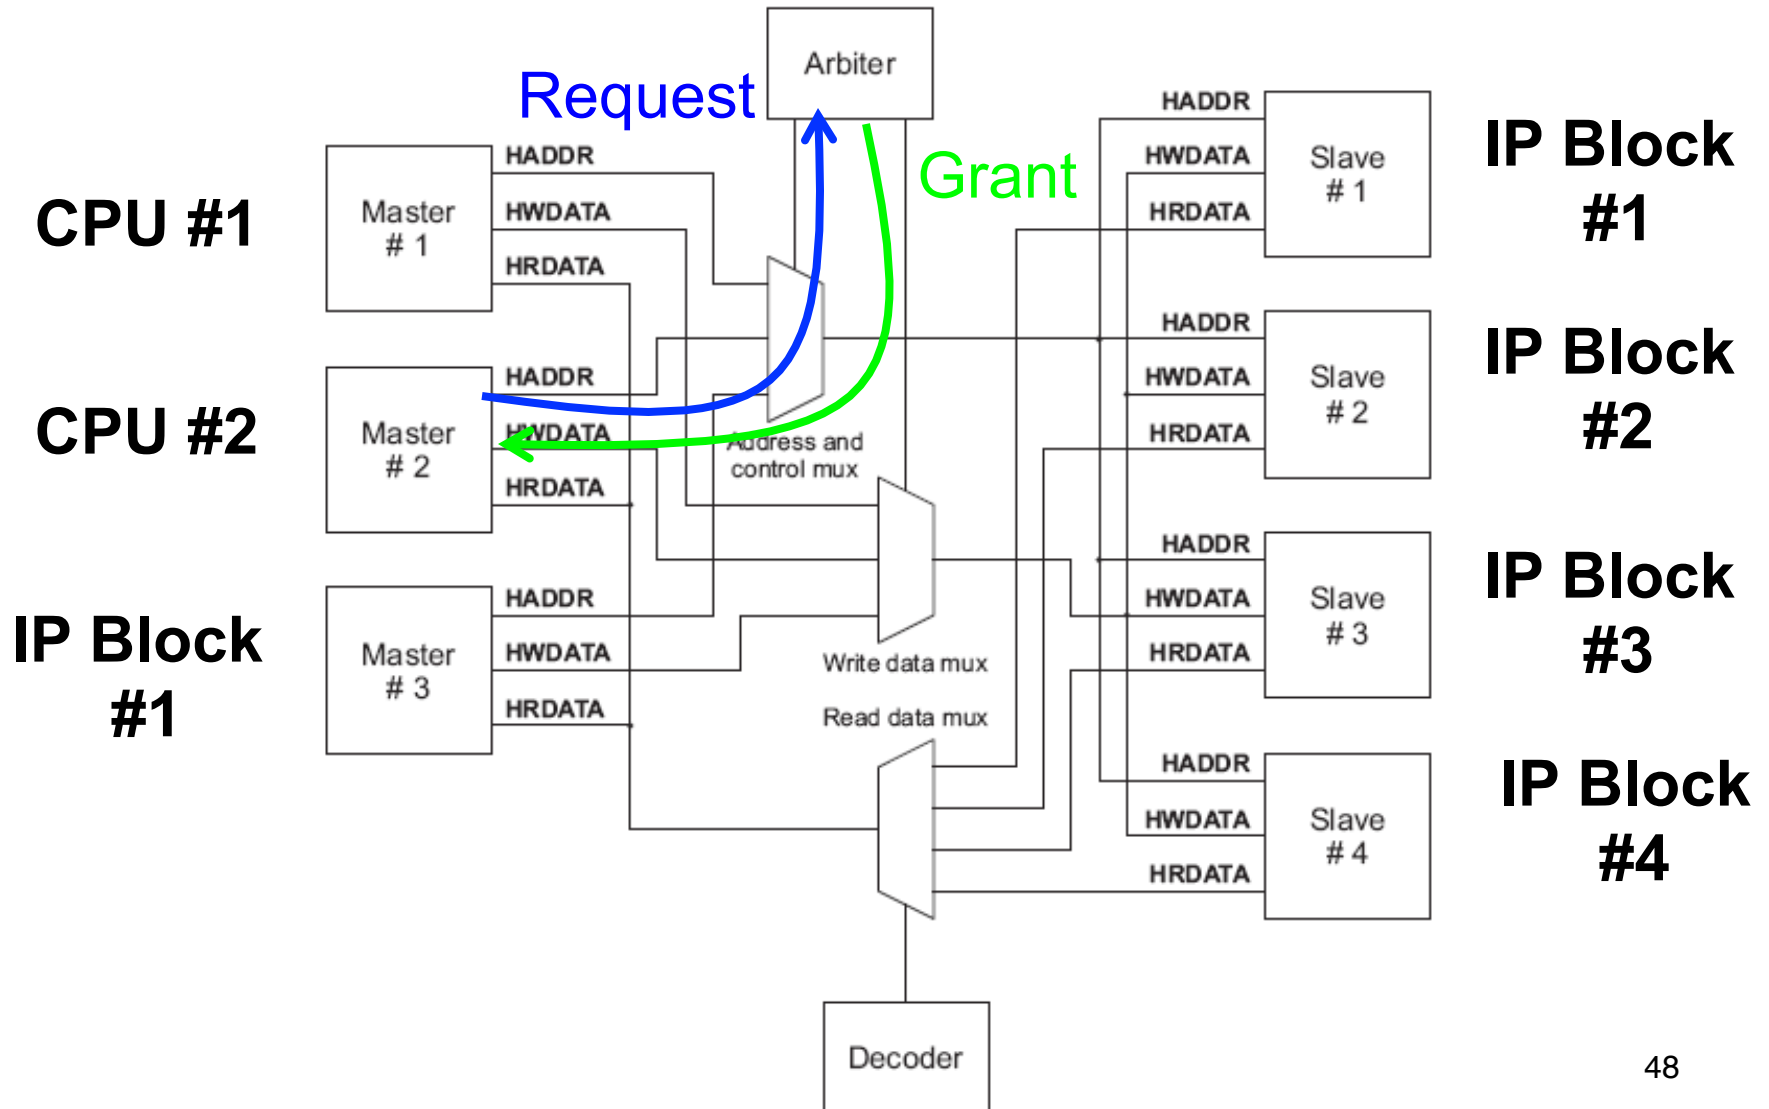
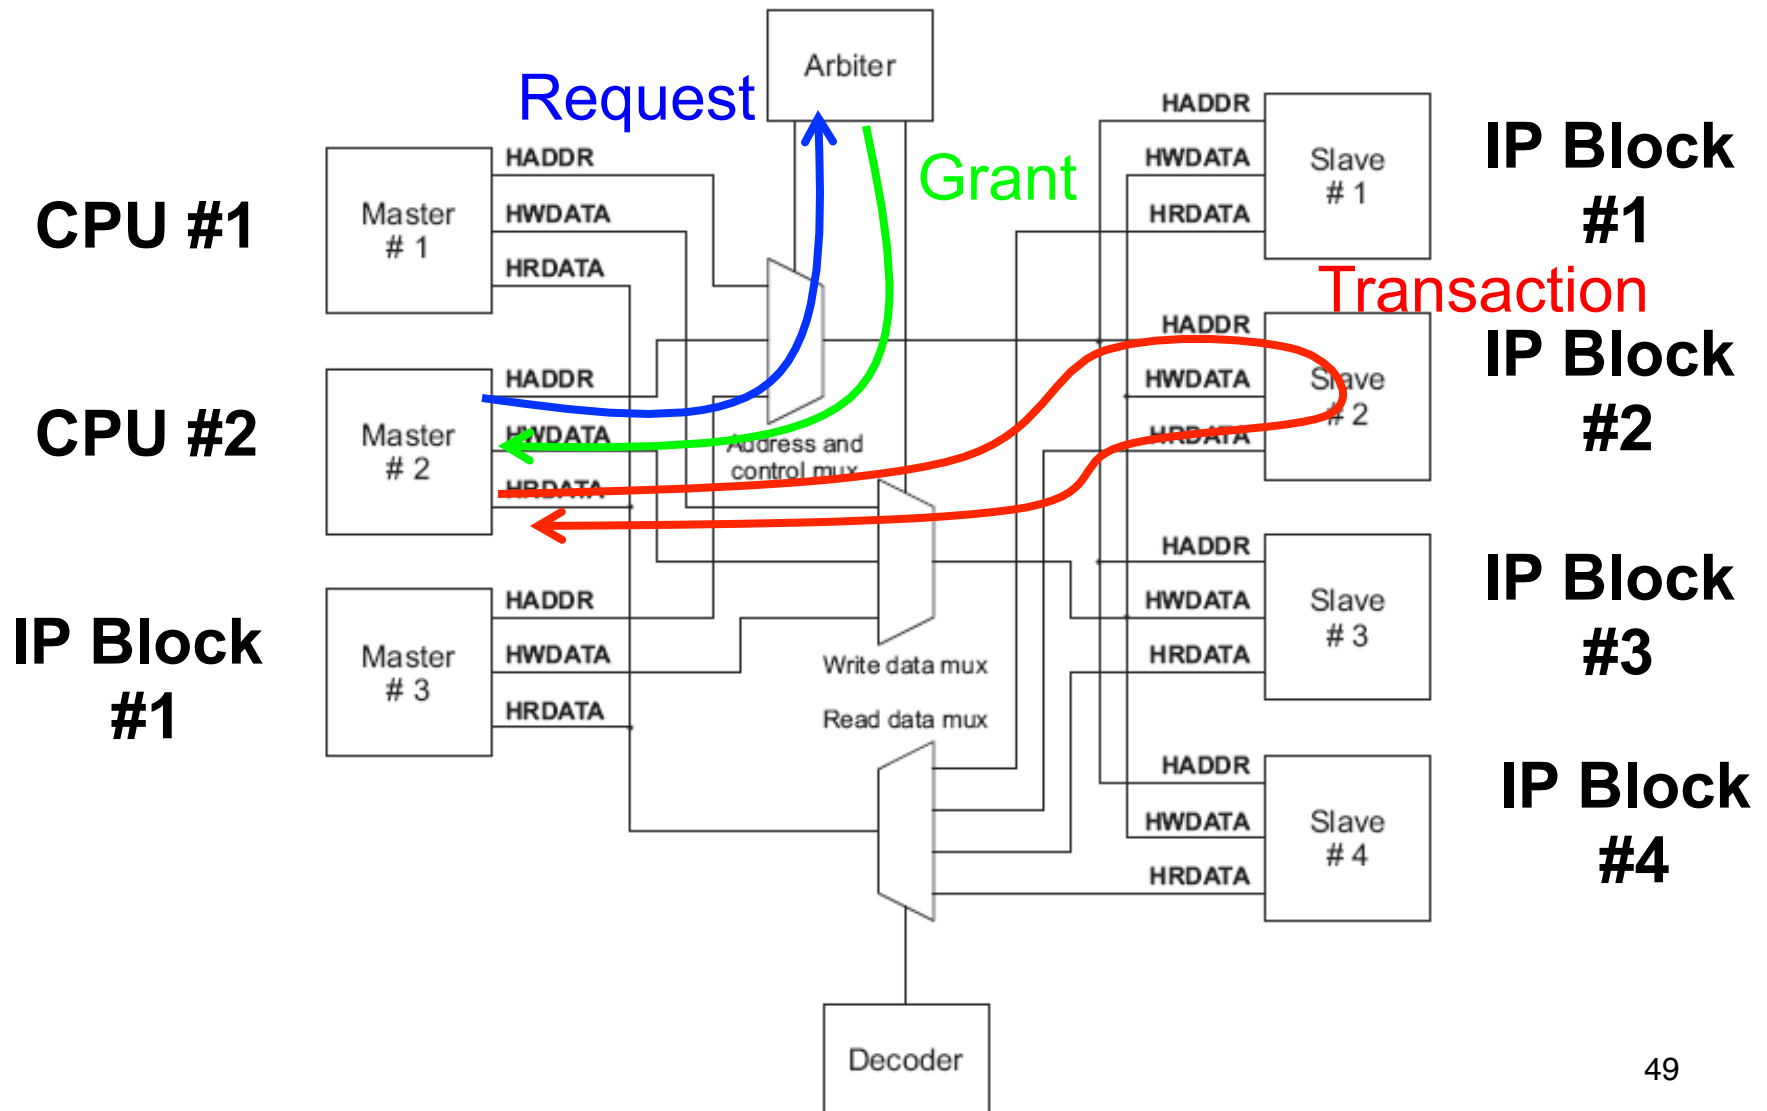# Bring on the complexity...
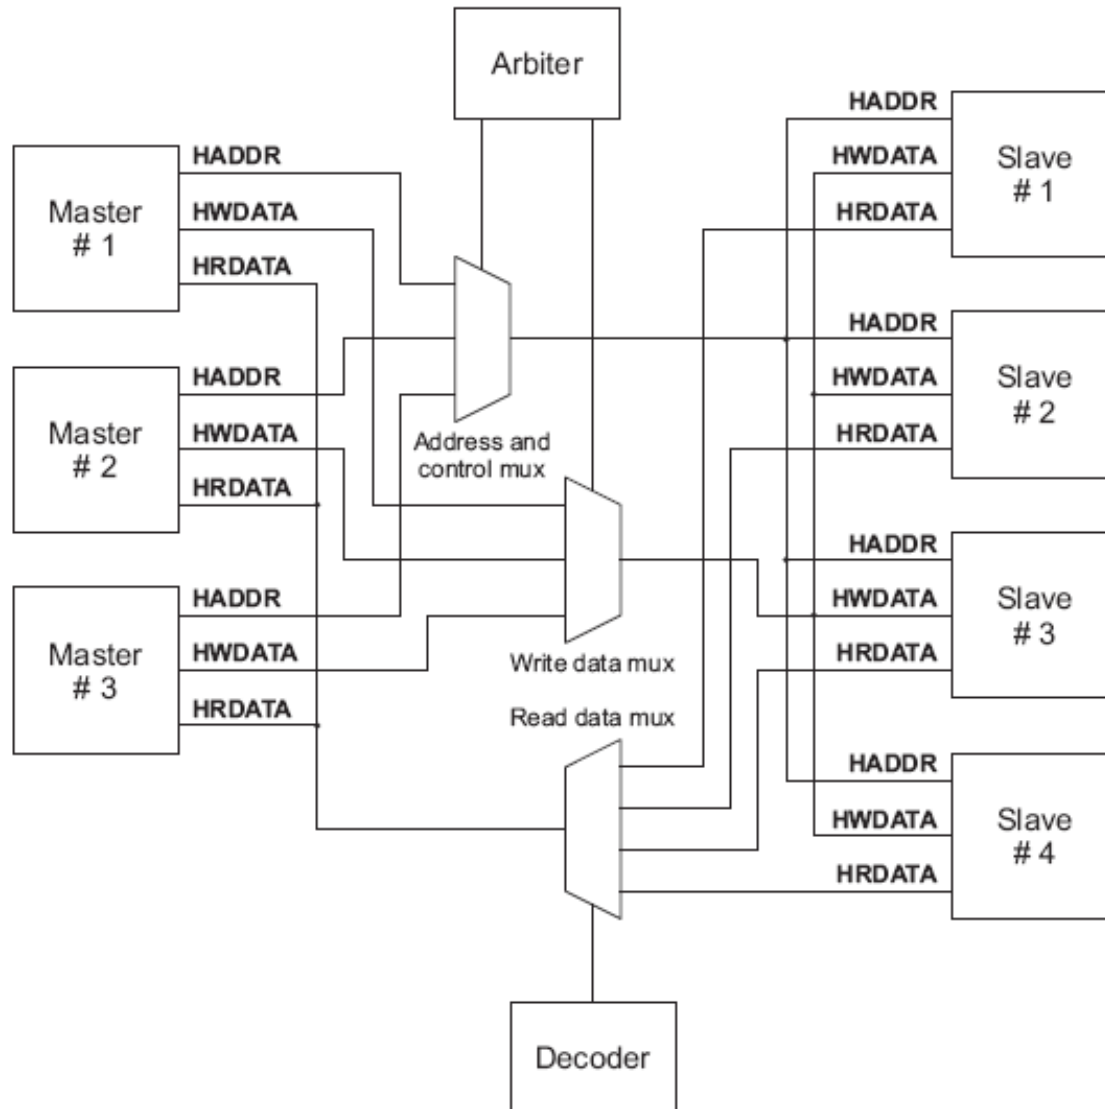
Bring on the complexity...

CPU #1

CPU #2

IP Block #1

IP Block #1

IP Block #2

IP Block #3

IP Block #4

Bring on the complexity...

Bring on the complexity...

Bring on the complexity...

Arbiter

Request

Grant

Transaction

CPU #1

CPU #2

IP Block #1

Master #1

HADDR
HWDATA
HRDATA

Master #2

HADDR
HWDATA
HRDATA

Master #3

HADDR
HWDATA
HRDATA

Address and control mux

Write data mux

Read data mux

Decoder

HADDR
HWDATA
HRDATA

Slave #1

HADDR
HWDATA
HRDATA

Slave #2

HADDR
HWDATA
HRDATA

Slave #3

HADDR
HWDATA
HRDATA

Slave #4

IP Block #1

IP Block #2

IP Block #3

IP Block #4

49
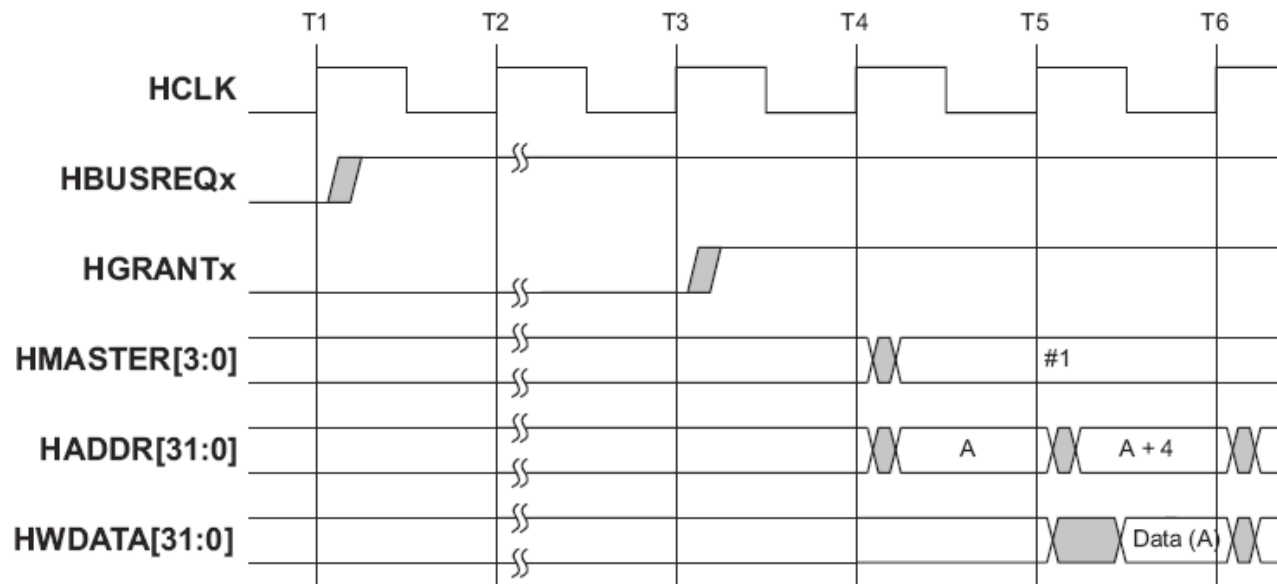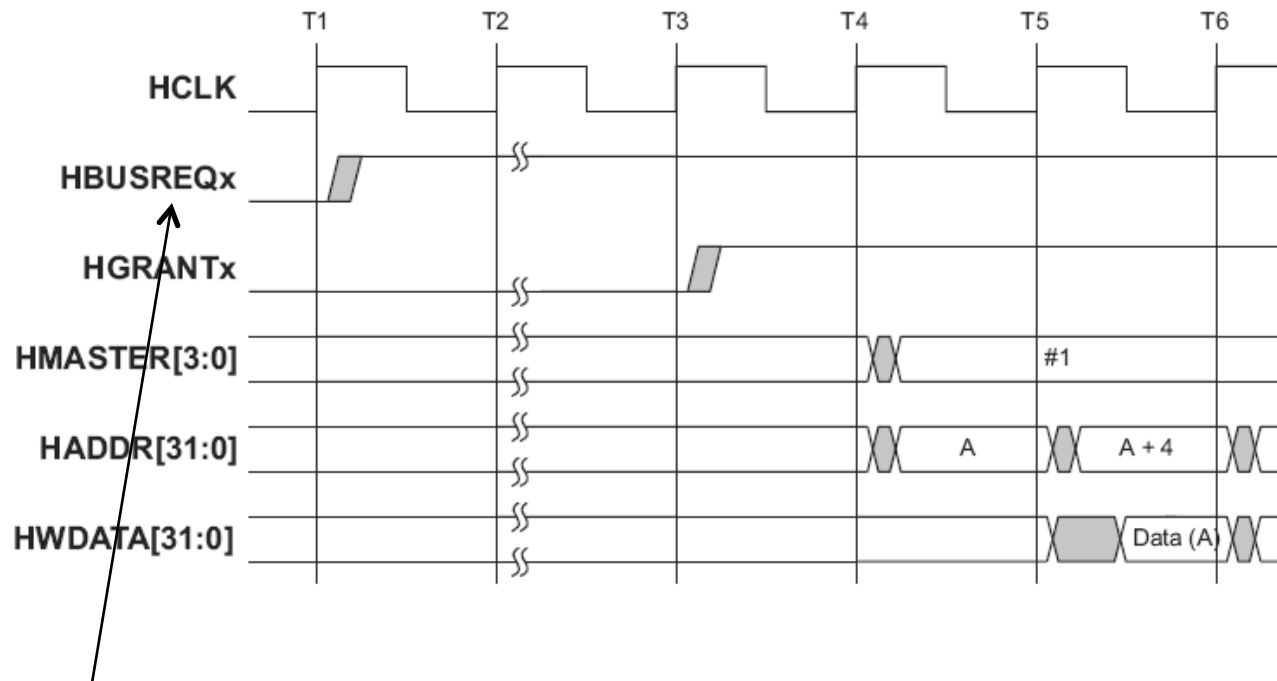
Bus Arbitration

- When multiple masters share a bus there must be some central resource to manage the bus: an **arbiter**

- Once there is competition for the bus, it is possible that it is not ready when you need it: **backpressure**

- Backpressure adds complexity and hurt performance
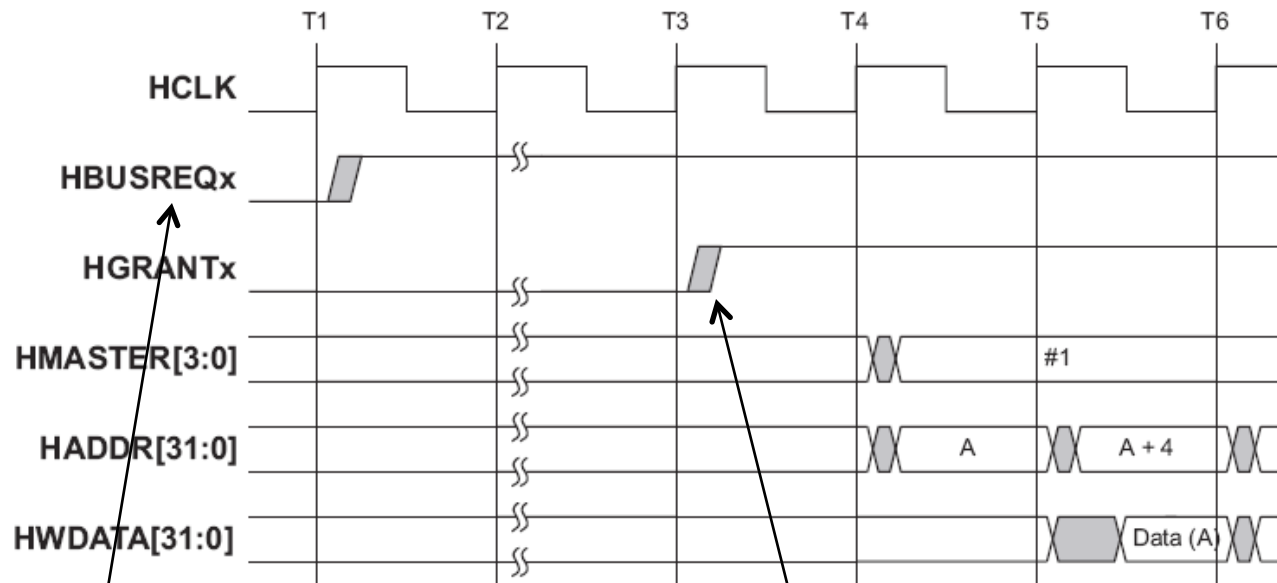
Request / Grant Protocol

Request / Grant Protocol

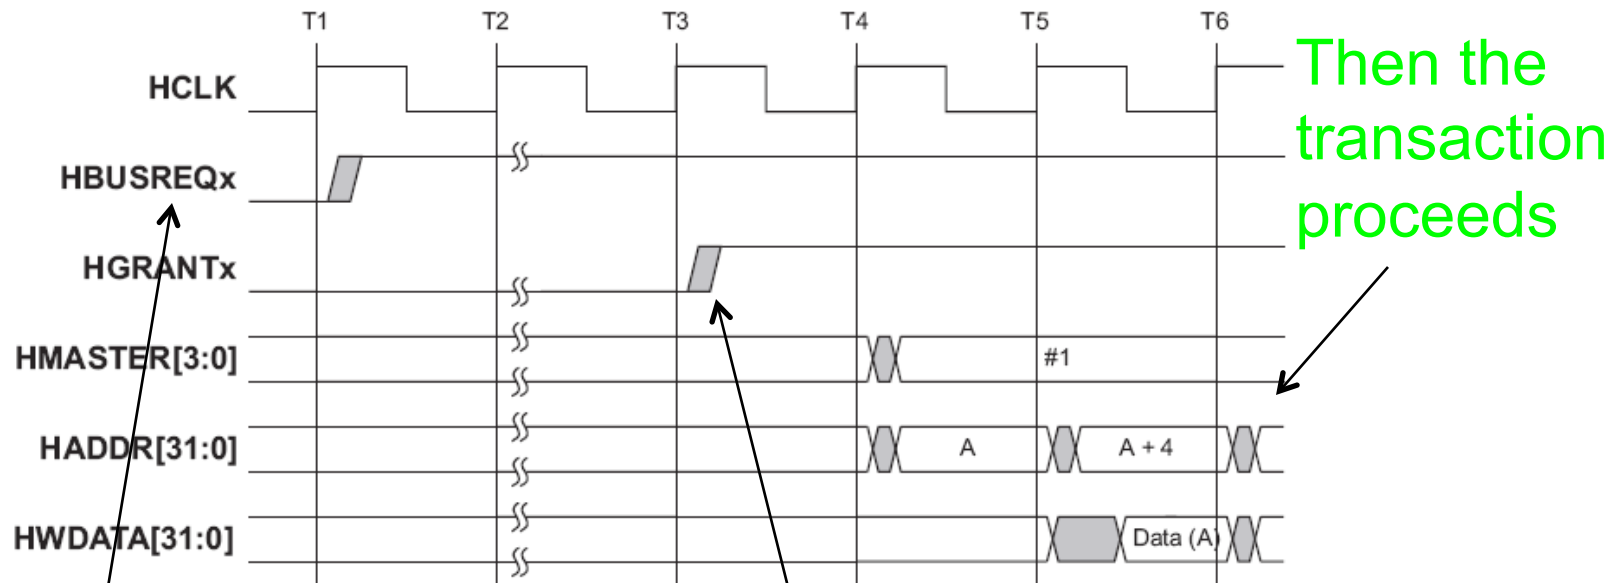Before a transaction a master makes a request to the central arbiter

Request / Grant Protocol

Before a transaction a master makes a request to the central arbiter

Eventually the request is granted

53

Request / Grant Protocol

Then the transaction proceeds

Before a transaction a master makes a request to the central arbiter

Eventually the request is granted

Request / Grant Protocol

Performance Impact

Then the transaction proceeds

Before a transaction a master makes a request to the central arbiter

Eventually the request is granted

Pipelined Transactions

- To help improve **bus efficiency** the transactions on the bus can be pipelined

- This is really a simple implementation of **multiple outstanding transactions**

- The address for one transaction can be presented before the data from the **previous transaction** has been completed
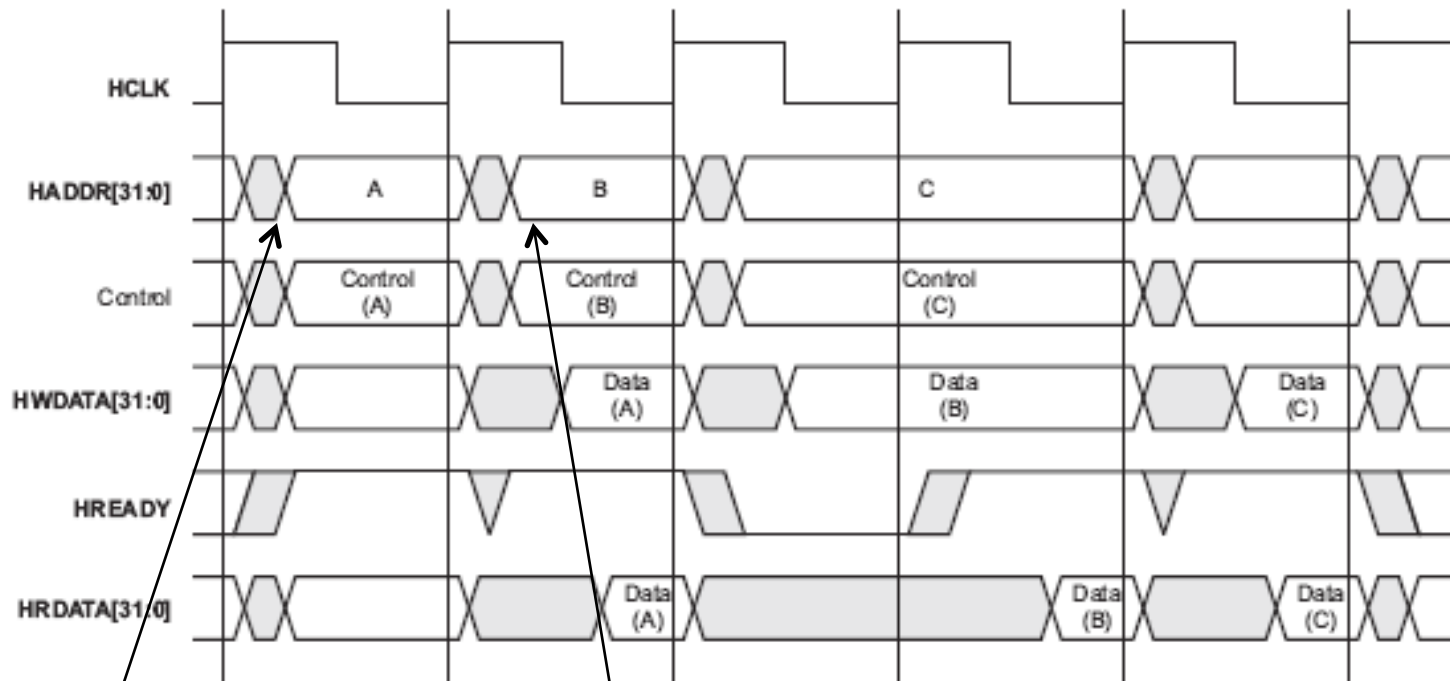
Pipelined Transactions
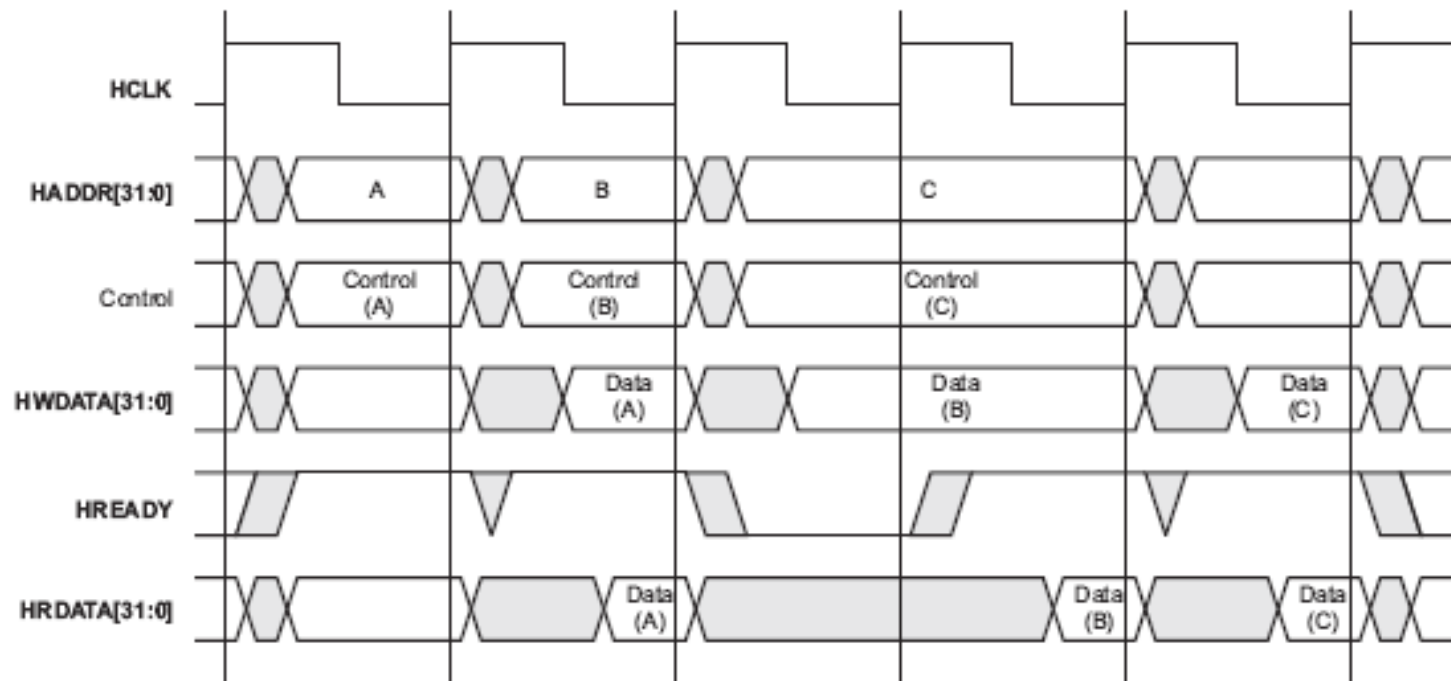
Pipelined Transactions
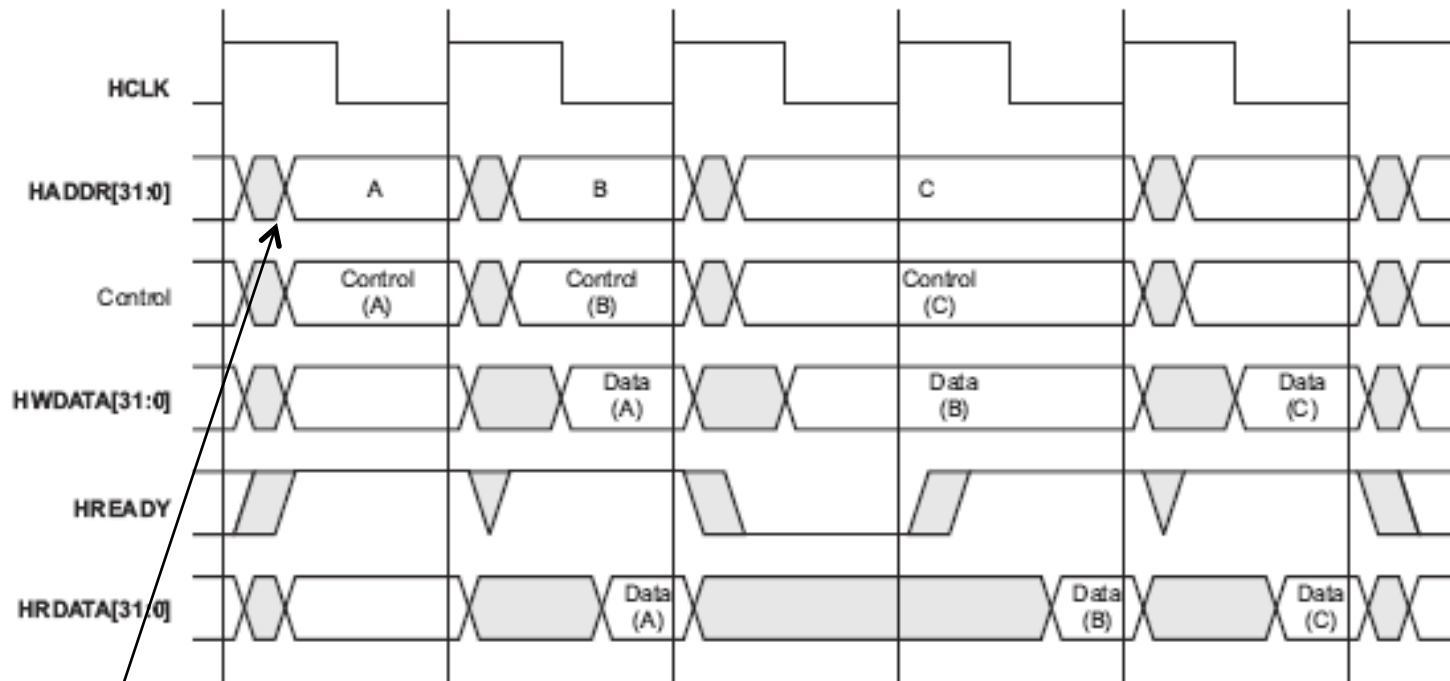
Transaction A Starts

Pipelined Transactions

Transaction A Starts

Transaction B Starts

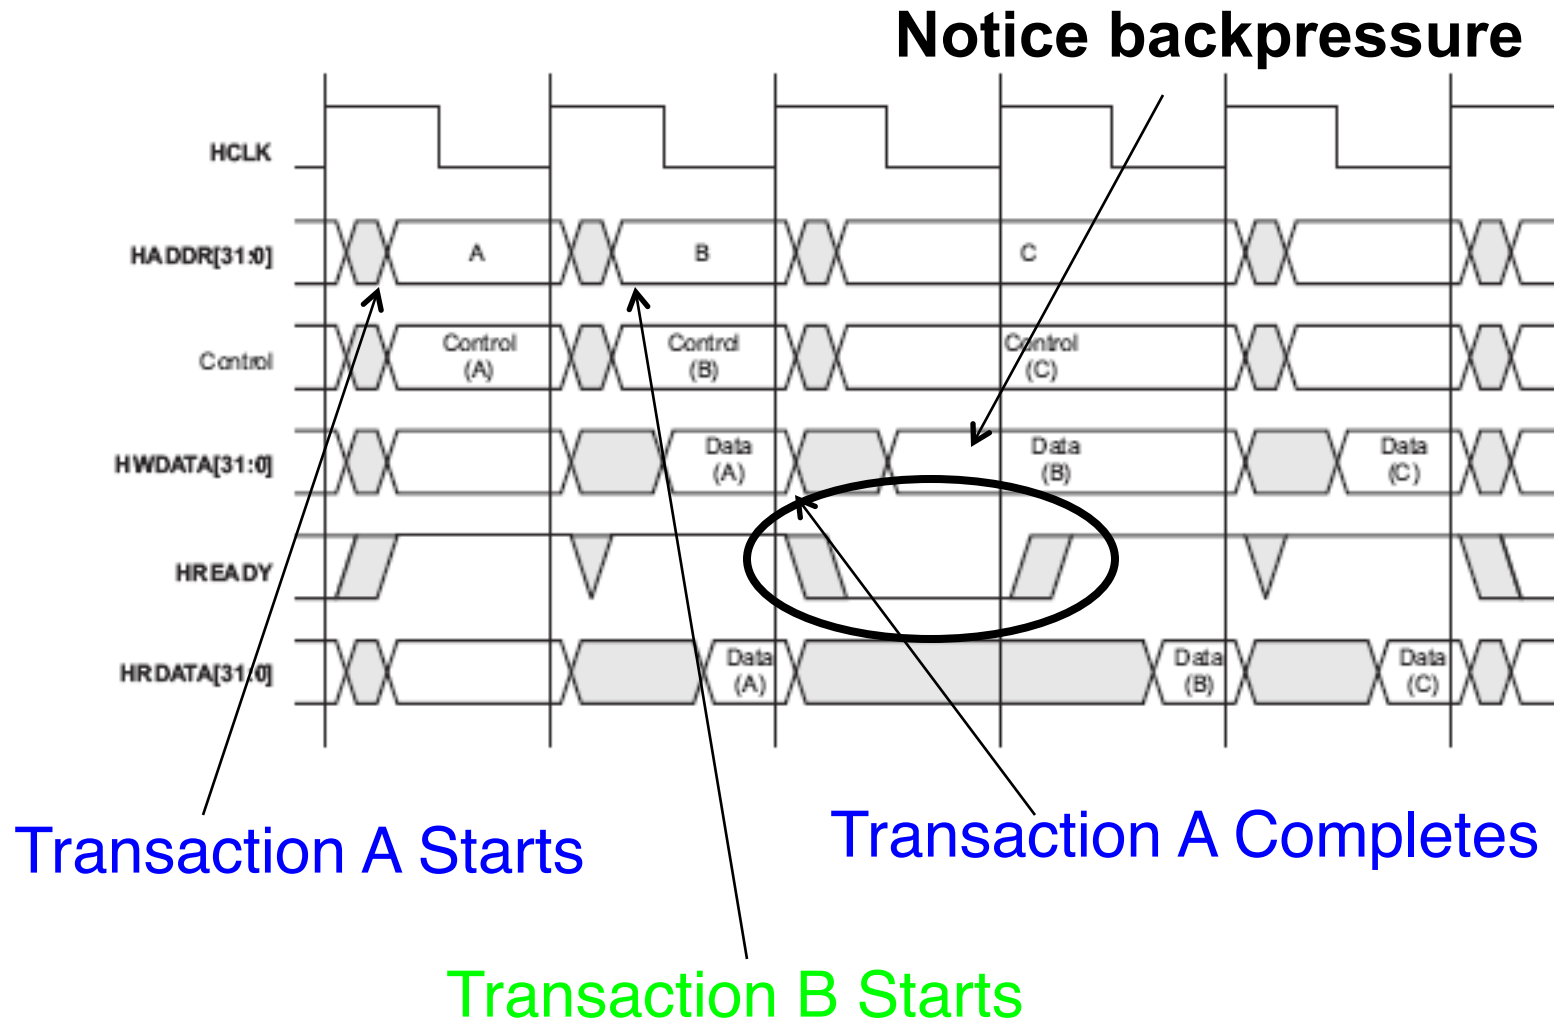# Pipelined Transactions

Pipelined Transactions

Notice backpressure

Transaction A Starts

Transaction A Completes
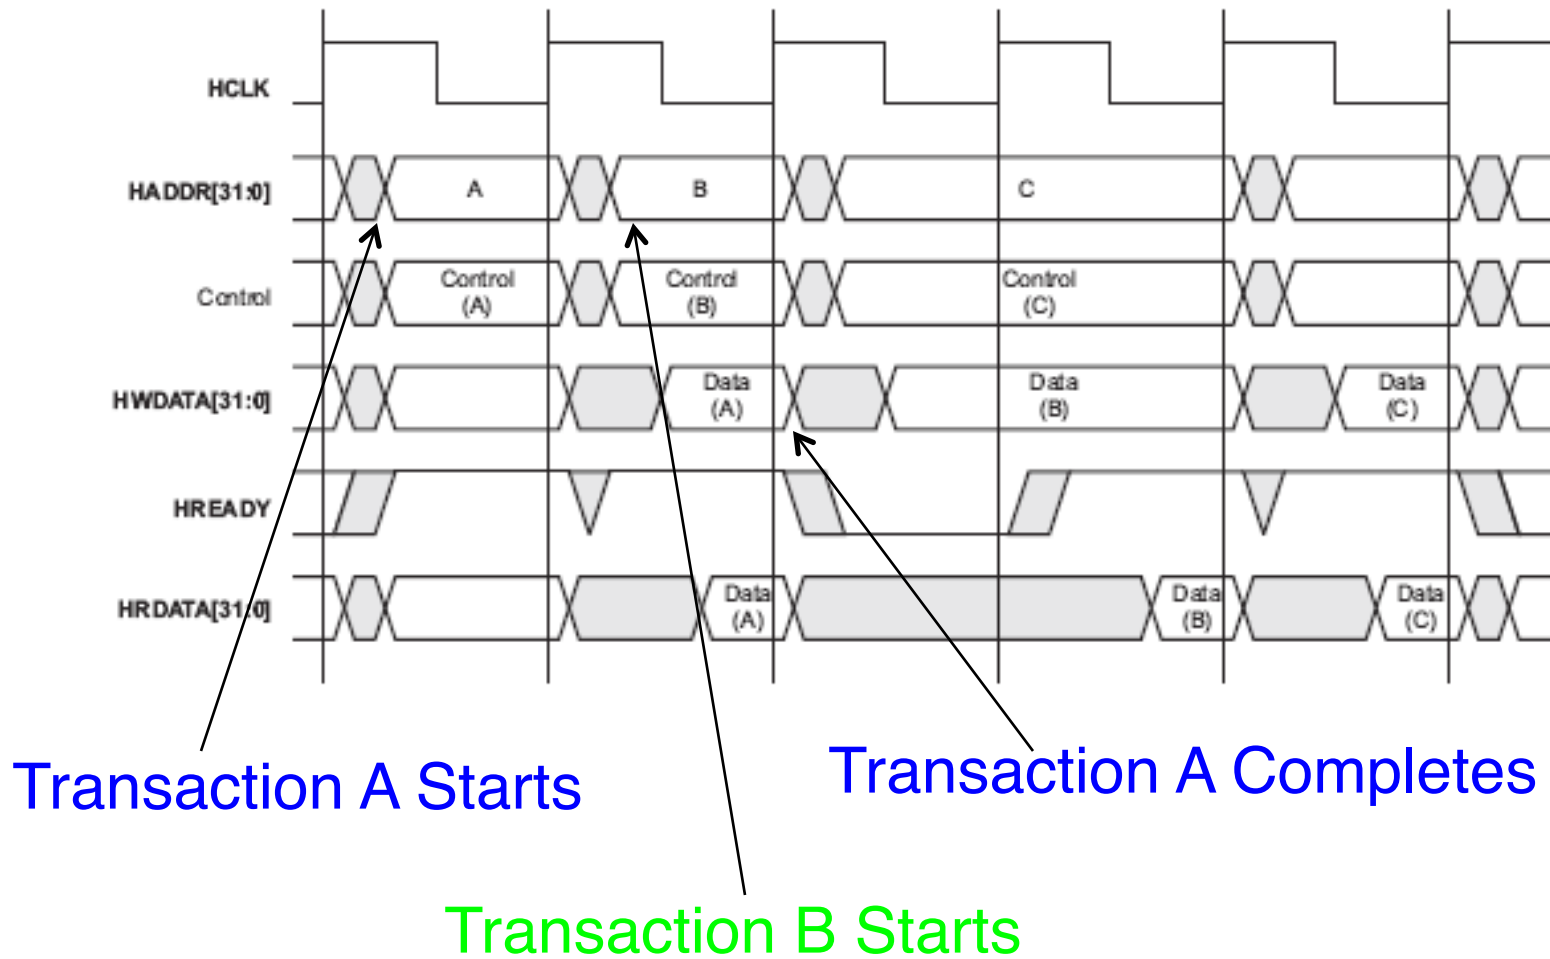
Transaction B Starts

Advantages

- Relatively easy to add new blocks
- Still has the familiar bus structure
- Low hardware cost
- Bus arbitration "solves" many ordering problems

Disadvantages

- Busses that require arbitration:
    - must route signals to the arbitration logic *and back*
    - must find a "fair" way to share the bus
    - slaves are not always available => backpressure
    - difficult to provide performance guarantees...

- Still potentially a bandwidth bottleneck

- Still doesn't scale well when blocks are added

- Multiple outstanding transactions not handled well - no ordering information

Networks-on-Chip (NoCs)

Networks-on-Chip

- It is clear that even with significant design effort the bus-style interconnect is not going to sufficient for large SoCs:

  - the **physical implementation** does not scale: bus fanout, loading, arbitration depth all reduce operating frequency

  - the available **bandwidth** does not scale: the single bus must be shared by all masters and slaves

Networks-on-Chip

- It is clear that even with significant design effort the bus-style interconnect is not going to sufficient for large SoCs:

  – the **physical implementation** does not scale: bus fanout, loading, arbitration depth all reduce operating frequency

  – the available **bandwidth** does not scale: the single bus must be shared by all masters and slaves

- **Lets start again:** Leverage research from data networking

What do we want?

- The SoCs of the future will:

  - have 100s of hardware blocks,
  - have billions of transistors,
  - have multiple processors,
  - have large wire-to-gate delay ratios,
  - handle large amounts of high-speed data,
  - need to support "plug-and-play" IP blocks

- Our NoC needs to be ready for these SoCs...

The Ideal Network

- What would the ideal network look like?:

    - Low area overhead
    - Simple implementation
    - High-speed operation
    - Low-latency
    - High-bandwidth
    - Operate at a constant frequency even with additional blocks
    - Increase available bandwidth as blocks are added
    - Provide performance guarantees
    - Have a "universal" interface

The Ideal Network

- What would the ideal network look like?:

  - Low area overhead
  - Simple implementation
  - High-speed operation
  - Low-latency
  - High-bandwidth
  - Operate at a constant frequency even with additional blocks
  - Increase available bandwidth as blocks are added
  - Provide performance guarantees
  - Have a "universal" interface

These are competing requirements: Design a network that is the "best" fit.

What do we need to decide?

- Network Interface
- Network Protocol / Transaction Format
- Network Topology
- VLSI Implementation

Network Interface

- We want our network to be "plug-and-play" so **industry standardization** is key

- However the standard be **universal** enough to address many different needs

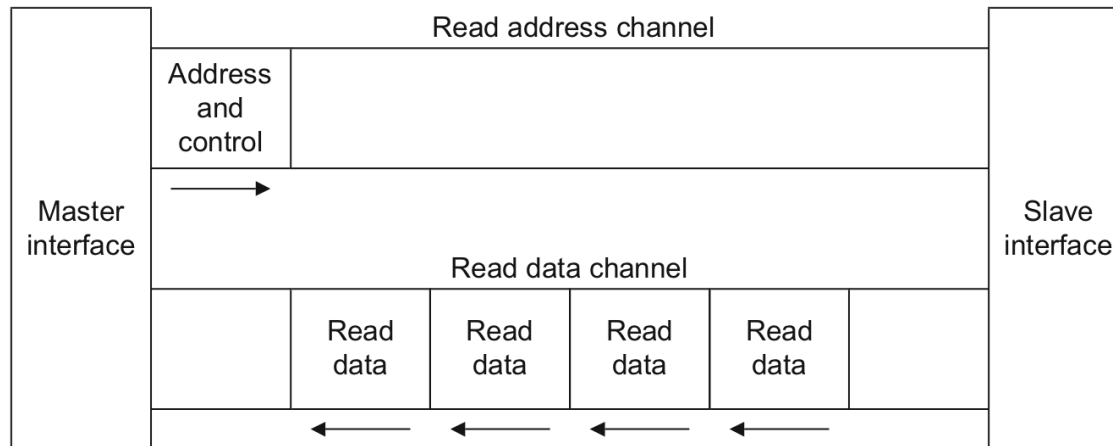- AMBA AXI is an example of an attempt at this

AMBA AXI

- ARM added the AXI specification to Version 3.0 of the AMBA standard

- **New approach:** define the interface and leave the interconnect up to the designers

- Good plan since a specific bus implementation is no longer required

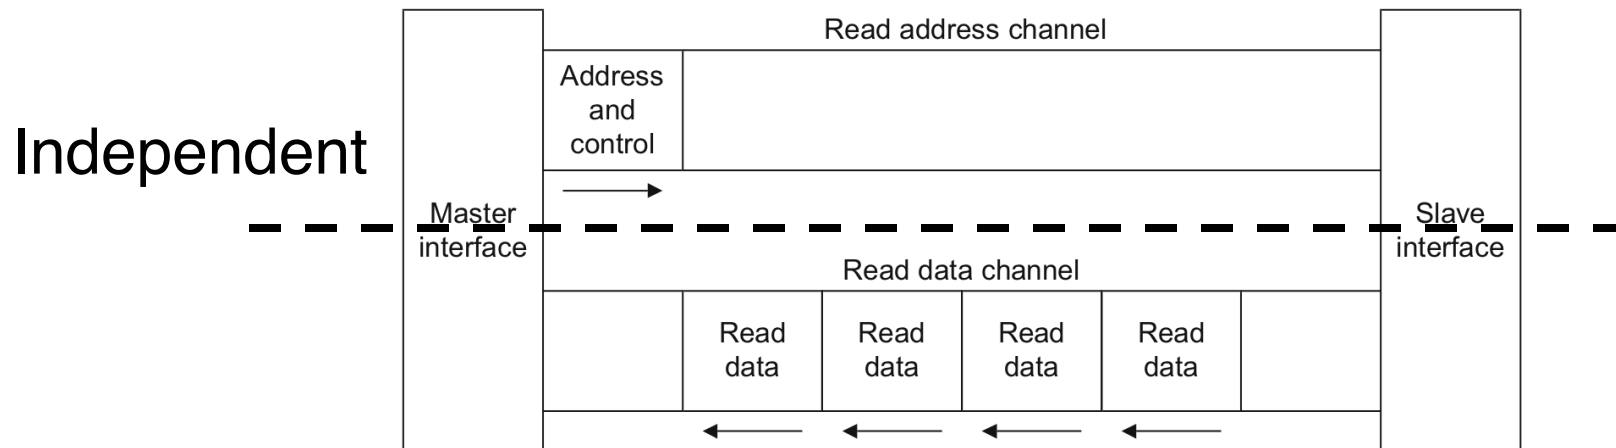- It is possible to use AXI to build many different NoCs

AMBA AXI

- Interface divided into 5 channels:

  – Write Address
  – Write Data
  – Write Response
  – Read Address
  – Read Data/Response

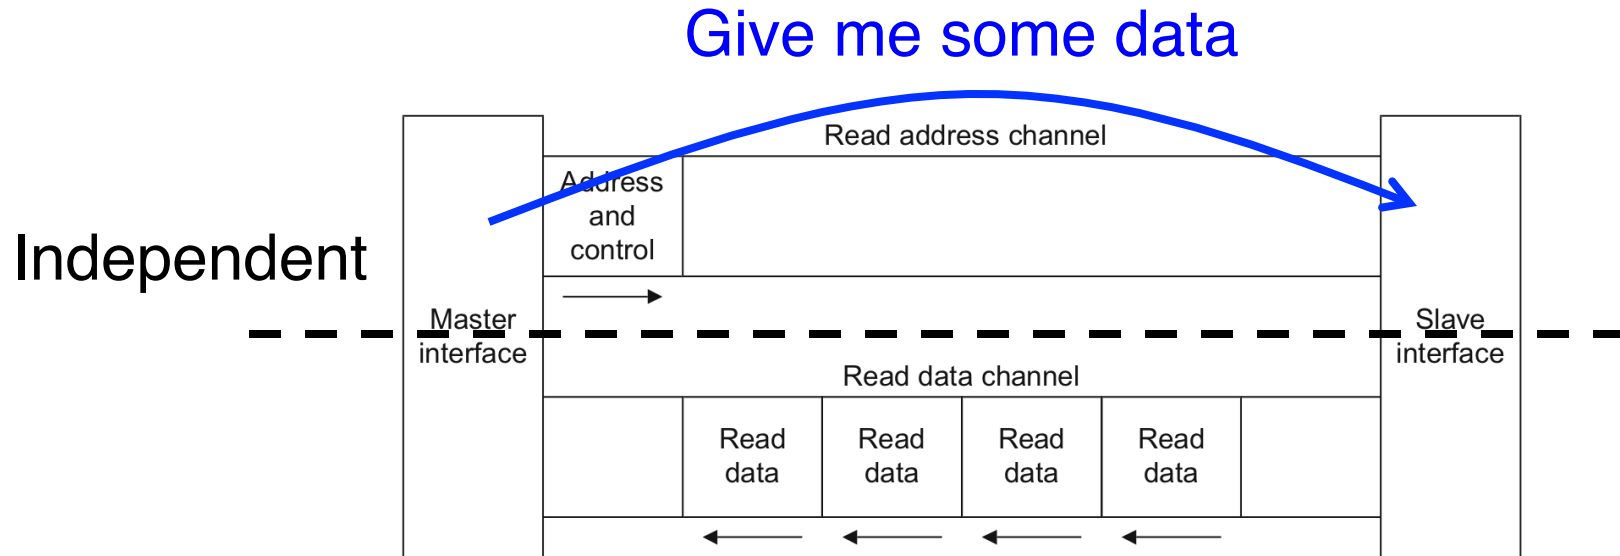- Each channel is **independent** and use **two-way flow control**
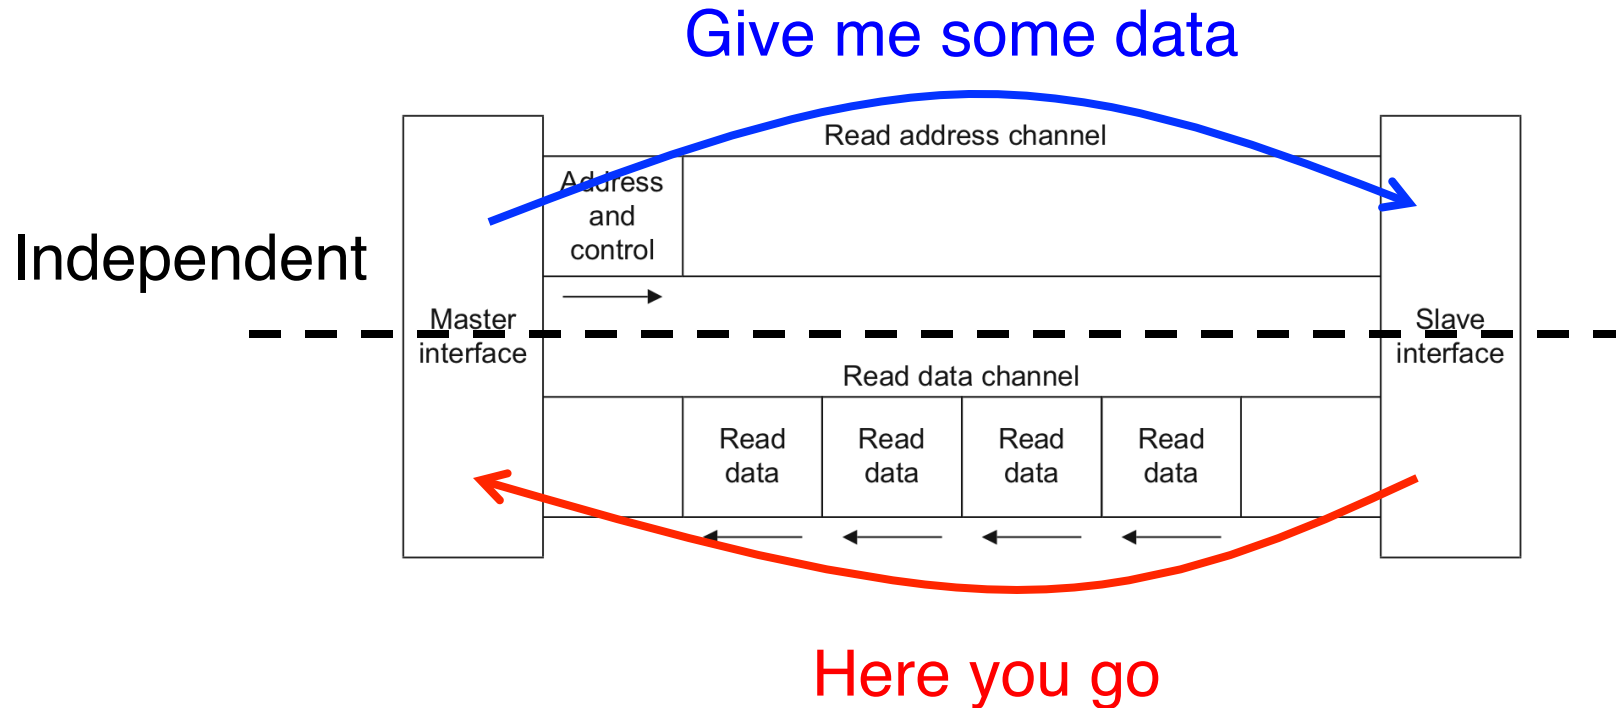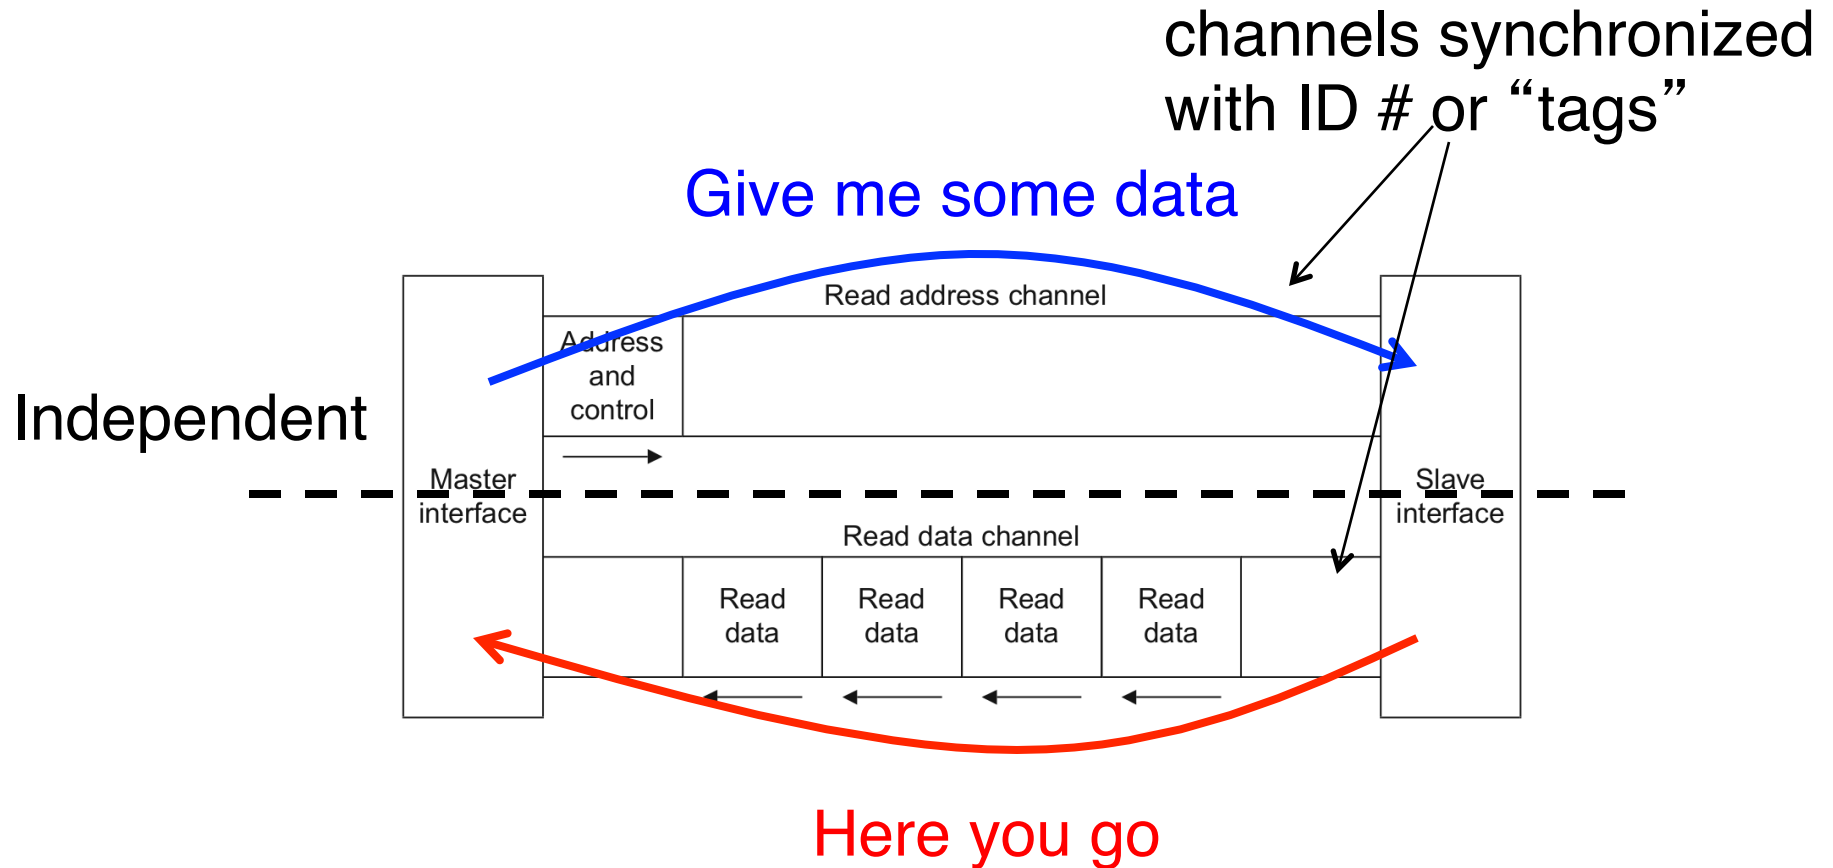
AMBA AXI Read Channels

AMBA AXI Read Channels

Independent

AMBA AXI Read Channels

Give me some data

Independent

Read address channel

Address
and
control

Master
interface

Slave
interface

Read data channel

| Read data | Read data | Read data | Read data |

AMBA AXI Read Channels

Give me some data

Independent

Read address channel

Address and control

Master interface

Slave interface

Read data channel

Read data	Read data	Read data	Read data

Here you go

AMBA AXI Read Channels

channels synchronized
with ID # or "tags"

Give me some data

Independent

Master
interface

Read address channel

Address
and
control

Slave
interface

Read data channel

| Read data | Read data | Read data | Read data | |

Here you go

AMBA AXI Write Channels

AMBA AXI Write Channels

Independent

Independent

AMBA AXI Write Channels

I'm sending data. Please store it.

Independent

Independent

AMBA AXI Write Channels

AMBA AXI Write Channels

83

AMBA AXI Write Channels

I'm sending data. Please store it.

Independent

Independent

Here is the data.

I received that data correctly.
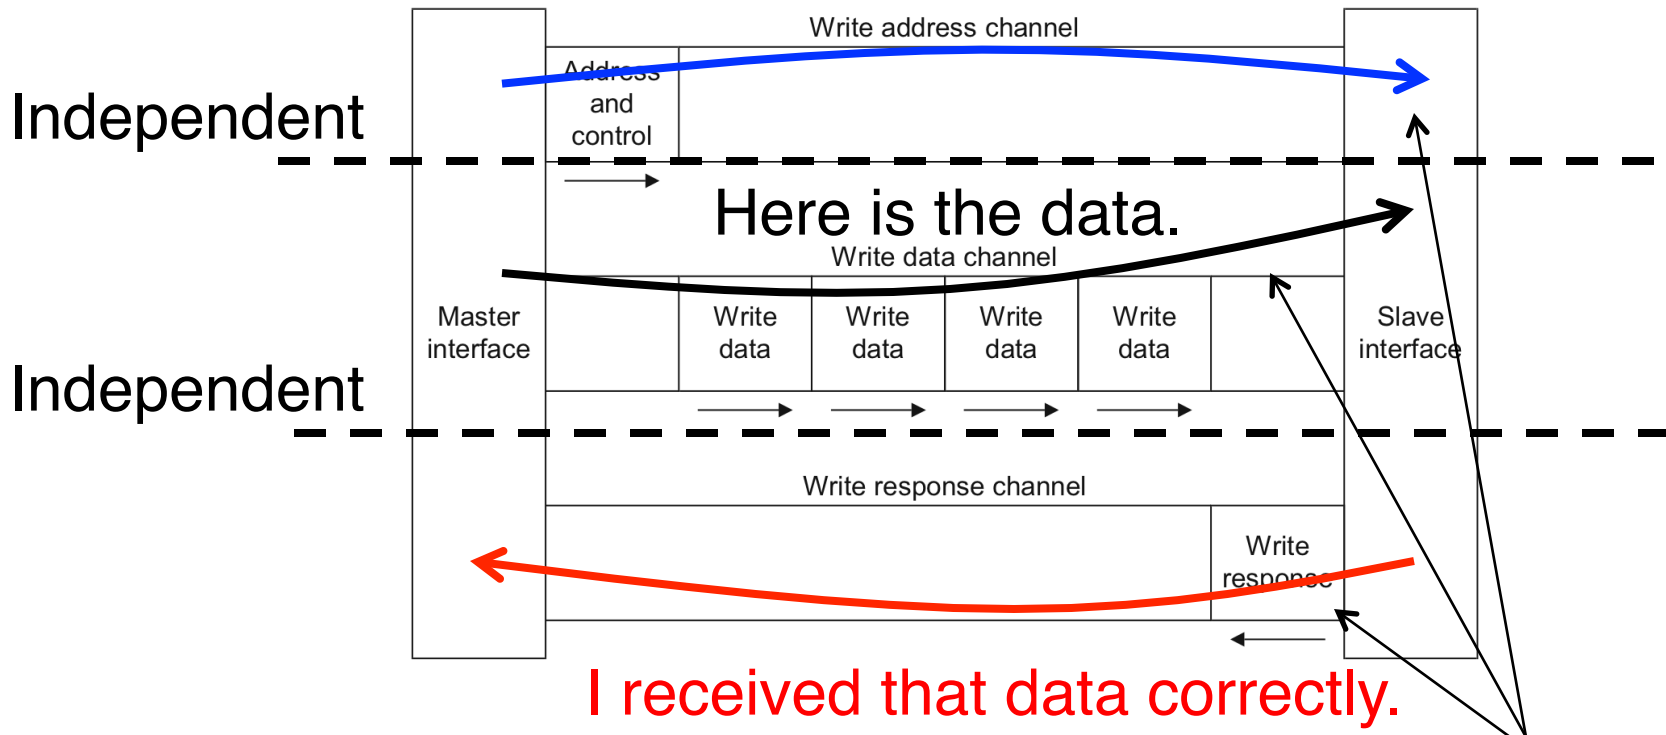
channels synchronized
with ID # or "tags"

AMBA AXI Flow-Control

- **Information moves only when:**

  - Source is **Valid**, and
  - Destination is **Ready**

- **On each channel the master or slave can limit the flow**

- **Very flexible**

AMBA AXI Flow-Control

- Information moves only when:

  – Source is **Valid**, and
  – Destination is **Ready**

- On each channel the master or slave can limit the flow

- Very flexible

Transfer

AMBA AXI Flow-Control

- This definition of very independent, fully flow-controlled channels is very useful

- However, there is a potential problem:

AMBA AXI Flow-Control

- This definition of very independent, fully flow-controlled channels is very useful

- However, there is a potential problem: **DEADLOCK**

AMBA AXI Flow-Control

- This definition of very independent, fully flow-controlled channels is very useful
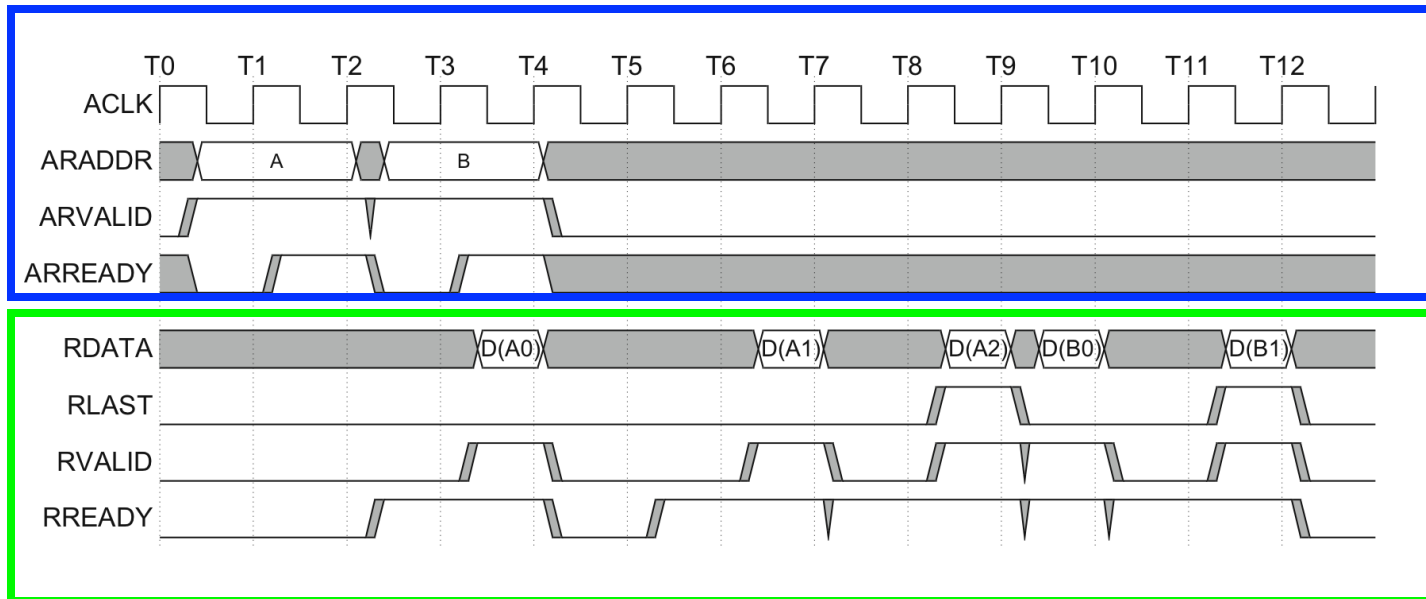
- However, there is a potential problem: **DEADLOCK**

- On a write transaction the master **must not wait** for AWREADY before asserting WVALID
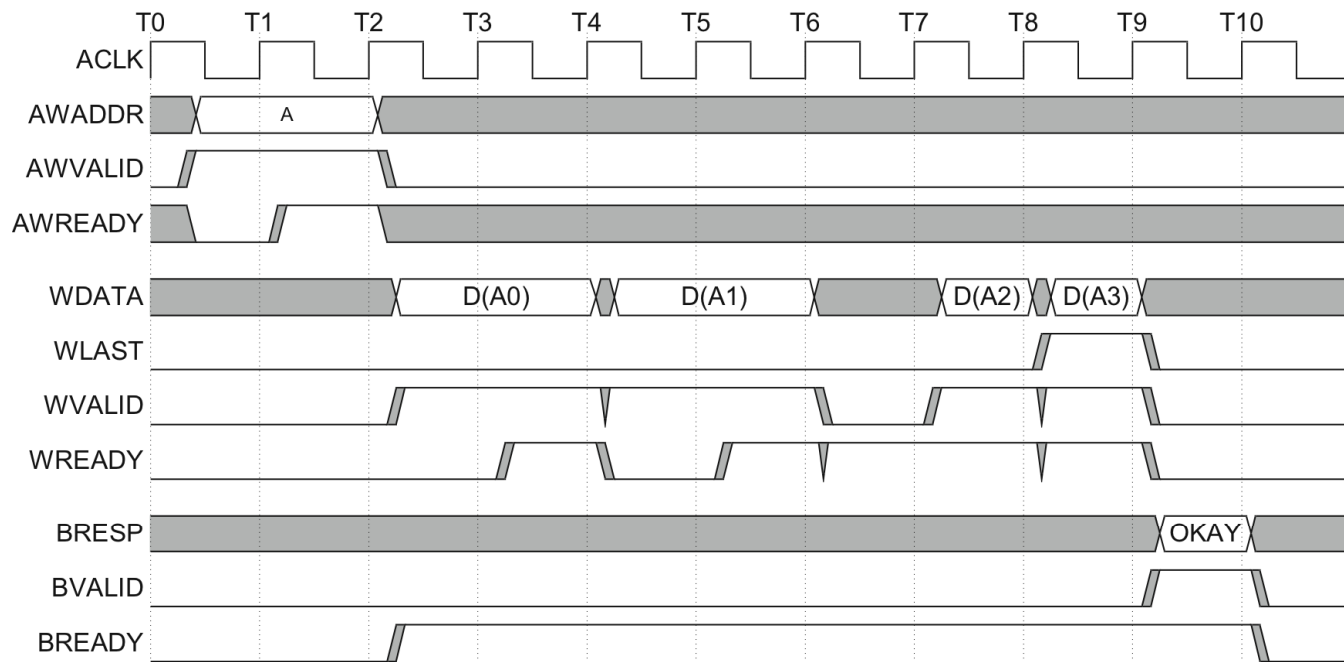
AMBA AXI Read

AMBA AXI Read

Read Address Channel

Read Data Channel

91

AMBA AXI Write

AMBA AXI Write

Write Address Channel

Write Data Channel

Write Response Channel

A True Interface Specification

- Because of the **channel independence** and the two-way **flow-control** the interface does not dictate the network protocol, transaction format, network topology, or VLSI implementation

- For example:
  - if you want to build a packet-based network, you can "backpressure" the data channel while you build the packet header from the address channel information,
  - you can use store-and-forward, or cut-through,
  - etc.

Network Protocol / Transaction Format

- There are many choice for network protocols and transactions formats:

    - **circuit-switched** : plan and provision a connection before communication starts

    - **packet-switched** : issues packets which compete for network resources

    - **hybrids:** schedule connectivity (dynamic or static)

Network Protocol / Transaction Format

- There are many choice for network protocols and transactions formats:

    - **circuit-switched** : plan and provision a connection before communication starts

    - **packet-switched** : issues packets which compete for network resources

    - **hybrids:** schedule connectivity (dynamic or static)

- There is still lots of research here....

Network Topology

- How should your network elements be interconnected:

  - **Fully Connected (N^2):** high area cost, high performance
  - **Mesh:** low area cost, potential poor performance
  - **Hypercube:** medium area, traffic dependent performance
  - **Fat-tree:** medium area, traffic dependent performance
  - **Torus:** medium area, traffic dependent performance

Network Topology

(a) SPIN   (b) CLICHE   (c) TORUS   (d) FOLDED TORUS   (e) IRREGULAR   (f) BFT

- There is lots of research here....

Network Topology - Caveat

- There has been a lot of research on topologies for NoCs, however it is important to realize that the performance of a topology is highly dependent on the **traffic patterns**!

- Traffic patterns in an SoC that you are designing yourself are **NOT** random, therefore much of the topology research is not applicable to most SoCs!
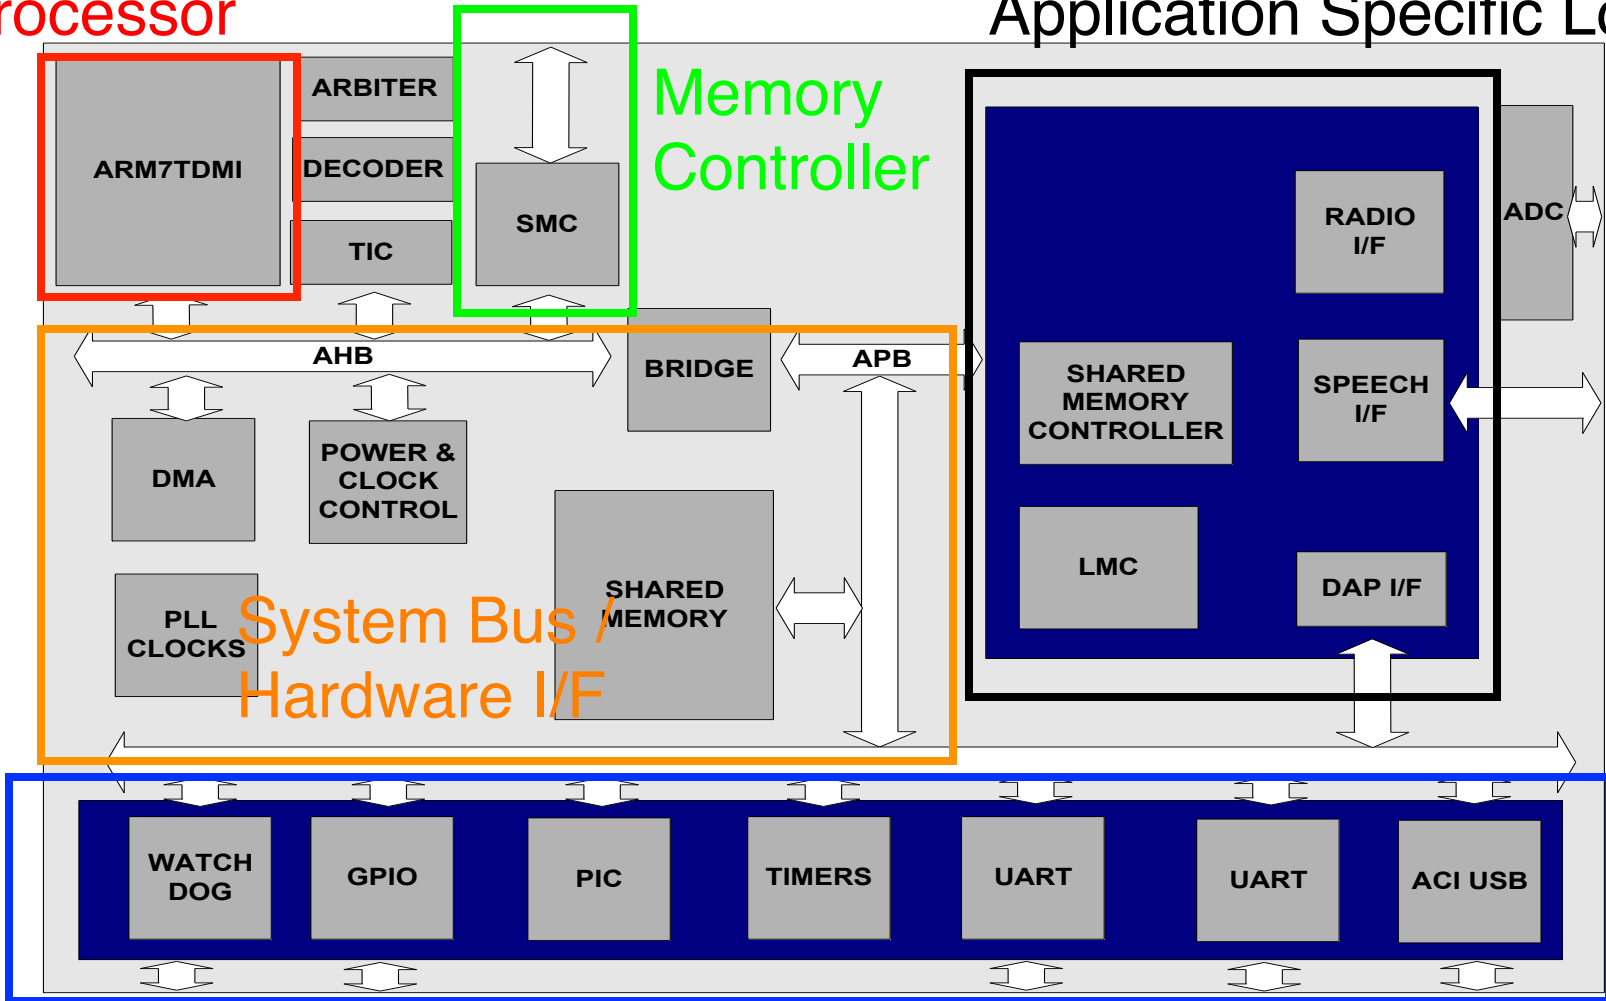
VLSI Implementation

- Once you have a topology there is still the mater of implementing it on your SoC

- There are many considerations:

  - **Clocking:** Synchronous, Asynchronous
  - **Buffer Insertion:** Trade-off power, area, performance
  - **Register Insertion / Pipelining:** Trade-off clock frequency, area, and latency
  - **Packet Buffers:** Trade-off area, latency and throughput

- Again, lots of research on-going...

Bluetooth "Platform" SoC

Processor

Application Specific Logic

ARBITER

DECODER

TIC

ARM7TDMI

Memory Controller

SMC

RADIO I/F

ADC

AHB

BRIDGE

APB

SHARED MEMORY CONTROLLER

SPEECH I/F

DMA

POWER & CLOCK CONTROL

PLL CLOCKS

System Bus / Hardware I/F

SHARED MEMORY

LMC

DAP I/F

WATCH DOG

GPIO

PIC

TIMERS

UART

UART

ACI USB

Low-speed I/O and Support Logic

101

Research Paper

- Lets look at:

Guerrier, P.; Greiner, A., "A generic architecture for on-chip packet-switched interconnections ," *Design, Automation and Test in Europe Conference and Exhibition 2000. Proceedings* , vol., no., pp.250-256, 2000